

Student Number

The University of Melbourne
Department of Computing and Information Systems

Final Examination, Semester 1, 2015
COMP10001 Foundations of Computing

Reading Time: 15 minutes. **Writing Time:** 2 hours.

This paper has 20 pages including this cover page.

Instructions to Invigilators:

Students must write all of their answers on this examination paper. Students may not remove any part of the examination paper from the examination room.

Instructions to Students:

There are 11 questions in the exam worth a total of 120 marks, making up 50% of the total assessment for the subject.

- All questions should be answered by writing a brief response or explanation in the lined spaces provided on the examination paper.
- It is not a requirement that all the lined spaces be completely filled; answers should be kept concise.
- Only material written in the lined spaces provided will be marked.
- The reverse side of any page may be used for notes or draft answers.
- Your writing should be clear; illegible answers will not be marked.
- Extra space is provided at the end of the paper for overflow answers.
Please indicate in the question you are answering if you use the extra space.
- Your answers should be based on Python 3 (the version that Grok uses), and can use any of the standard Python libraries.

Authorised Materials: No materials are authorised.

Calculators: Calculators are not permitted.

Library: This paper may NOT be held by the Baillieu Library.

<i>Examiners' use only</i>											
1	2	3	4	5	6	7	8	9	10	11	Total

Part 1: Algorithmic Thinking

Question 1

[10 marks]

Evaluate the following expressions, and provide the output in each case.

(a) `('comp', 10000 + 1)`

(b) `sorted({'dogs': ['andrew'], 'ducks': ['tim']})[-1]`

(c) `"c" + str(0)*2 + "lollipop"[:1][0]`

(d) `bool(range(2, 10, 2) and "dig" in "don't wait")`

(e) `[i for i in "how much wood could a woodchuck chuck".split() if len(i) > 6]`

Question 2

[8 marks]

Rewrite the following function, replacing the `while` loop with a `for` loop, but preserving the remainder of the original code structure:

```
def all_alpha(word):
    i = 0
    while i < len(word):
        if not word[i].isalpha():
            return False
        i += 1
    return True
```

[illegible]

Question 4**[12 marks]**

The following function `valid_url(url)` that takes a single string argument `url`, and is intended to return `True` if `url` is a valid full-specified URL (with the hostname ending in `.com` or `.au`), and `False` otherwise. Recall that a valid fully-specified URL is made up of a scheme, hostname, optional port and optional path.

```
def is_valid_url(url):
    SCHEMES = ("http", "ftp")
    TLDS = ("com", "au")
    try:
        scheme, rest = url.split("://")
    except ValueError:
        return False
    if scheme not in SCHEMES:
        return False
    rest_path = rest.split("/")
    hostname = rest_path[0]
    path = rest_path[1:]
    if ":" in hostname:
        (hostname, port) = hostname.split(":")
        try:
            port = int(port)
        except ValueError:
            return False
    tld = hostname.split(".")[1]
    if tld not in TLDS:
        return False
    return True
```

The provided implementation has logic errors in it, and rejects some valid fully-specified URLs it should accept, and conversely, accepts some inputs which are not valid fully-specified URLs.

- (a) Provide an example of a valid fully-specified URL for which `valid_url()` would return `True` (i.e. the function would perform correctly).

- (b) Provide an example of a valid fully-specified URL for which `valid_url()` would return `False`, despite it being valid.

- (c) Provide an example of an *invalid* fully-specified URL for which `valid_url()` would return `True`, despite it being invalid.

- (d) Provide an example of an *invalid* fully-specified URL for which `valid_url()` would return `False` (i.e. the function would perform correctly).

Part 2: Constructing Programs

Question 6

[10 marks]

The function `csv_aggregate(infile, outfile)` takes an input file name `infile` (a string) and output file name `outfile` (also a string). `infile` is of the following format:

```
docid,sentid,pos
0,0,0
0,1,1
1,0,0
1,1,0
1,2,0
```

where each (non-header) row contains (in order) the ID of a document, the ID of a sentence within that document, and a binary value indicating whether the given sentence has positive sentiment (i.e. says positive things) or not. The function takes `infile` and calculates the relative proportion of positive sentiment sentences in a given document, and returns a single row for that document, containing the document ID (`docid`) and the calculated proportion (`pos_ratio`) of positive sentiment sentences. This is saved in the file `outfile`. For example, given the file `sentences.csv` with the content above, `outfile` will be as follows:

```
docid,pos_ratio
0,0.5
1,0.0
```

Provide code to insert into each of the numbered boxes in the code on the next page to complete the function as described. Note that your code will be evaluated at the indentation level indicated for each box.


```
import csv

def csv_aggregate(infile, outfile):
    
    fpout.write("docid,pos_ratio\n")
    prev_docid = None
    
    for row in csv.DictReader(open(infile)):
        
        if prev_docid != None:
            fpout.write("{}{}\n".format(prev_docid,pos/total))
            
            pos = total = 0
        if row["pos"] == '1':
            pos += 1
            total += 1
    if prev_docid != None:
        
    fpout.close()
```

(1) _____

(2) _____

(3) _____

(4) _____

(5) _____

Question 7

[12 marks]

Write a function `equiword(word)` that takes a single argument `word` (a non-empty string) and returns a (positive) integer `n` if all unique letters in `word` occur exactly `n` times, and `False` otherwise.

For example:

```
>>> equiword("intestines")
2
>>> equiword("deeded")
3
>>> equiword("duck")
1
>>> equiword("doggy")
False
```

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and extend across the width of the page. There are no margins, text, or other markings on the paper.

Question 8**[18 marks]**

A “Takuzu” puzzle takes the form of an $n \times n$ square of even dimensions (e.g. 4×4 but not 5×5), each cell of which is to be filled with either a zero (0) or a one (1), such that, when solved:

- each row and column of the square contains an equal number of zeroes and ones; and
- no row or column contains more than 2 adjacent occurrences of the same digit.

For example, the following is a valid (completed) Takuzu square:

0	1	0	1
1	0	1	0
1	1	0	0
0	0	1	1

as is the following a valid (partially-solved) Takuzu square:

		0	1
1	0		
1		0	
	0		

whereas the following is not, because the first row and third column contain too many ones:

0	1	1	1
1		1	0
1	1	0	
0	0	1	1

and the following is similarly invalid, because it contains three adjacent ones in the first and fourth columns, and three adjacent zeroes in the third row:

0	1	0	1	1	0
1	0	1	0	1	0
1	1	0	0	0	1
1	0	0	1	1	0
0	0	1	1	0	1
0	1	0	1	0	1

We will represent Takuzu squares based on the following two-dimensional list representation (corresponding to the valid 4×4 solved example from the previous page), noting that the number of rows and columns will vary according to the size of the square:

```
takuzu = [[0, 1, 0, 1],
          [1, 0, 1, 0],
          [1, 1, 0, 0],
          [0, 0, 1, 1]]
```

For partially-solved Takuzu squares such as the following (which is valid):

		0	1
1	0		
1		0	
	0		

blank cells are represented by the value `None`, as follows:

```
takuzu = [[None, None, 0, 1],
          [1, 0, None, None],
          [1, None, 0, None],
          [None, 0, None, None]]
```

That is, there are three possible values for each cell: `None`, 0 and 1. Note that, for partially-solved Takuzu squares, the number of zeroes and ones in a given row or column may not be balanced, because of the presence of blank cells. Note also that, for the purposes of this question, an invalid square is one that violates one of the two constraints described above, and we are not concerned with the question of whether it has a unique solution or not, or the degree to which it is partially solved. As such, the following is a valid (partially-solved) square:

```
takuzu = [[None, None, None, None],
          [None, None, None, None],
          [None, None, None, None],
          [None, None, None, None]]
```

You may assume that all Takuzu inputs provided as arguments to your function are square, of even dimensions, at least of size 2×2 , contain one of the three possible values in each cell, and are either solved or partially solved.

Write a function `valid(takuzu)` that takes a Takuzu square as an argument and returns `True` if `takuzu` is a valid Takuzu square, and `False` otherwise.

For example:

```
>>> valid([[None, None, 0, 1], [1, 0, None, None], [1, None, 0, None],
... [None, 0, None, 1]])
True
>>> valid([[None, None, 0, 1], [1, 0, None, None], [1, None, 0, None],
... [1, 0, None, 1]])
False
>>> valid([[1, 0, 0, 1], [0, 0, 1, 1], [1, 1, 0, 0], [1, 0, 1, 0]])
False
>>> valid([[1, 0], [0, 1]])
True
>>> valid([[None, None], [None, None]])
True
>>> valid([[1, 1], [0, 1]])
False
```

This image shows a blank sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

[illegible]

Part 3: Conceptual Questions

Question 9: Algorithmic Problem Solving

[9 marks]

(a) In the context of algorithmic development, what is the difference between an “exact” and “approximate” approach?

[4 marks]

(b) Outline the “divide-and-conquer” strategy of algorithmic problem solving, with the aid of an example.

[5 marks]

Question 10: Applications of Computing**[10 marks]**

(a) With the aid of an example, describe what “alignment” means in the context of statistical machine translation, e.g. as used for language decipherment purposes.

[5 marks]

(b) List three (3) possible attacks on an “Internet voting system”.

[6 marks]

Question 11: HTML and the Internet**[12 marks]****(a)** Based on the following HTML document:

```
1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>HTML Example</title>
5 </head>
6 <body>
7 <ul>
8 <li>A <a href="http://en.wikipedia.org/wiki/Link">link</a>,</li>
9 <li>a link;</li>
10 <li>my kingdom for a link!</li>
11 </ul>
12 </body>
13 </html>
```

and the provided line numbers, identify in the document where each of the following items occurs. In the case that the item spans multiple lines, you should specify the full range of line numbers (e.g. 2–4).

[8 marks]

(i) A document type declaration:

(ii) A tag attribute:

(iii) Anchor text:

(iv) A list:

(b) What is the relationship between a “hostname” and an “IP address”?

[4 marks]

This is blank space for further answers should you need it. Please ensure that you label the answers in this area carefully, and that you indicate on the corresponding question page that your answer can be found here.

[illegible]

[illegible]

[illegible]

20 of 20