



INFO20003

Database Systems

David Eccles

Lecture 18

Database Administration



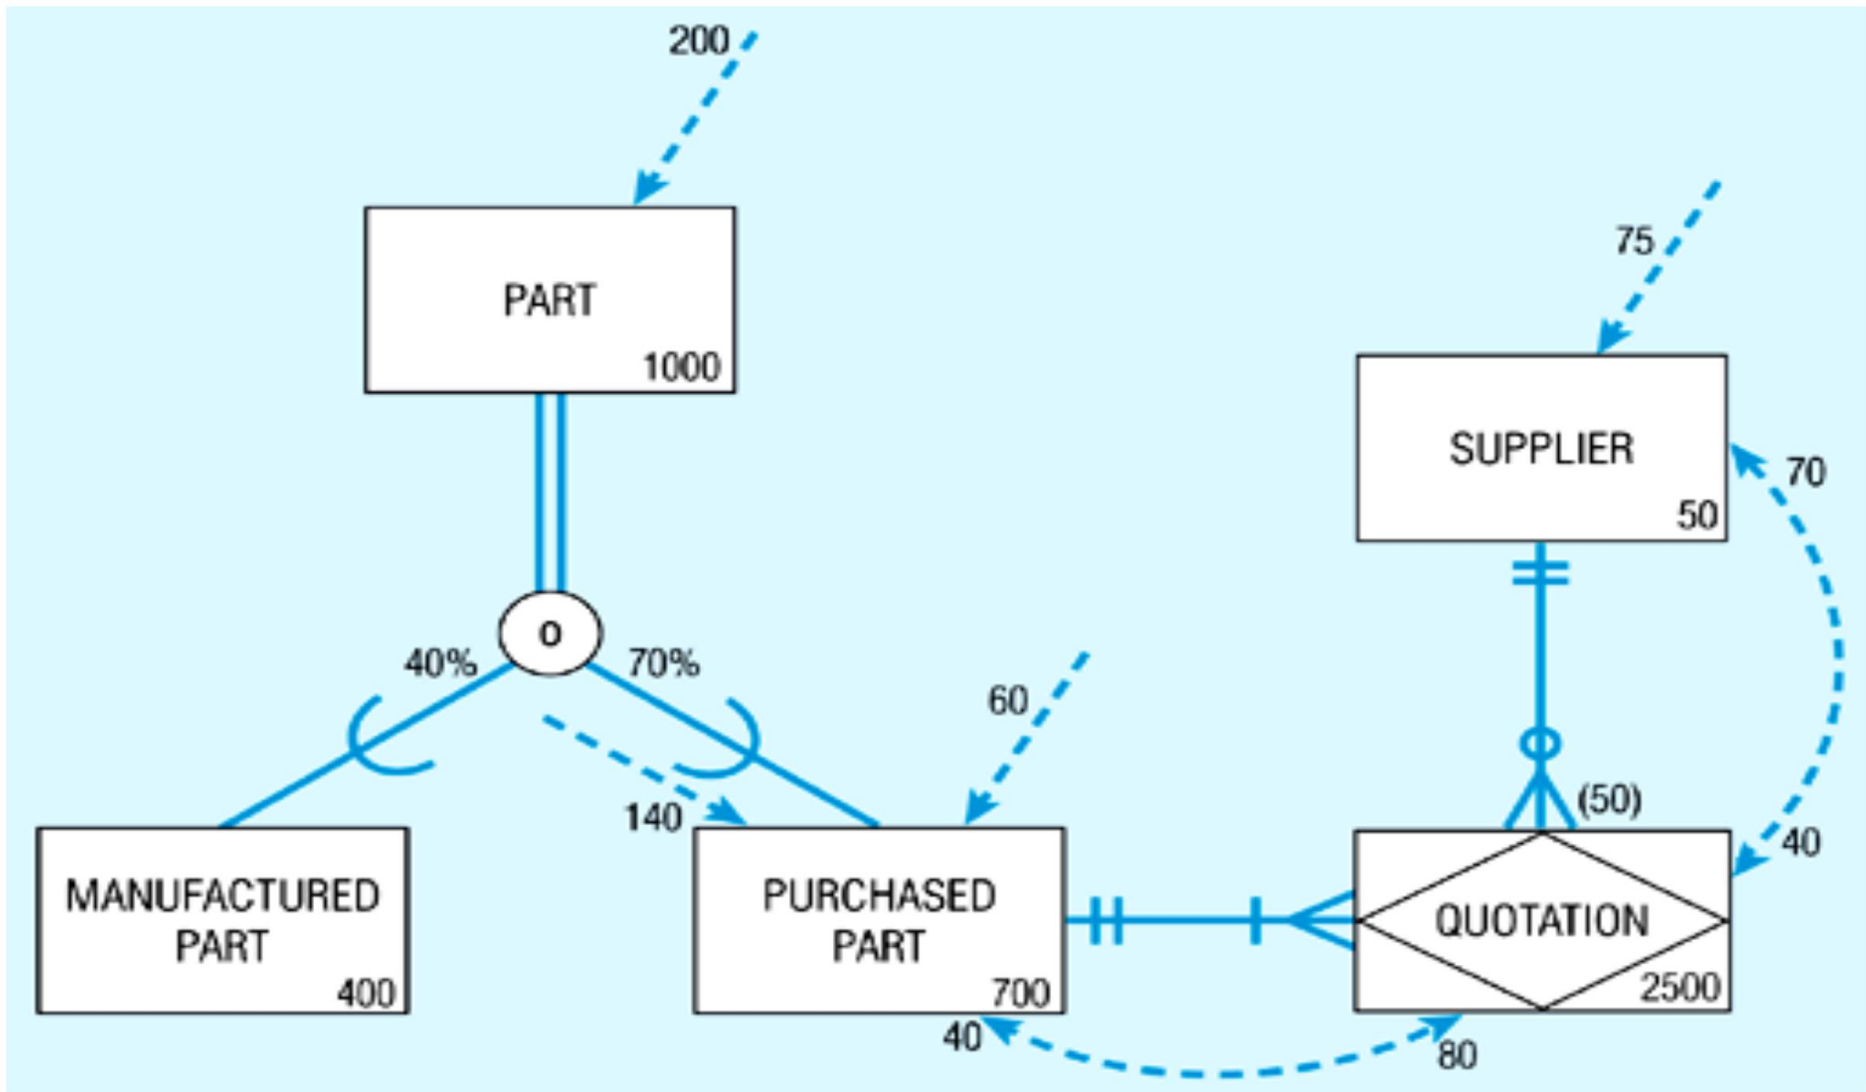
- Functions that are part of the DBA role
 - Capacity planning
 - Estimating disk space and transaction load
 - Backup and recovery
 - Types of failures, responses to these, types of backups

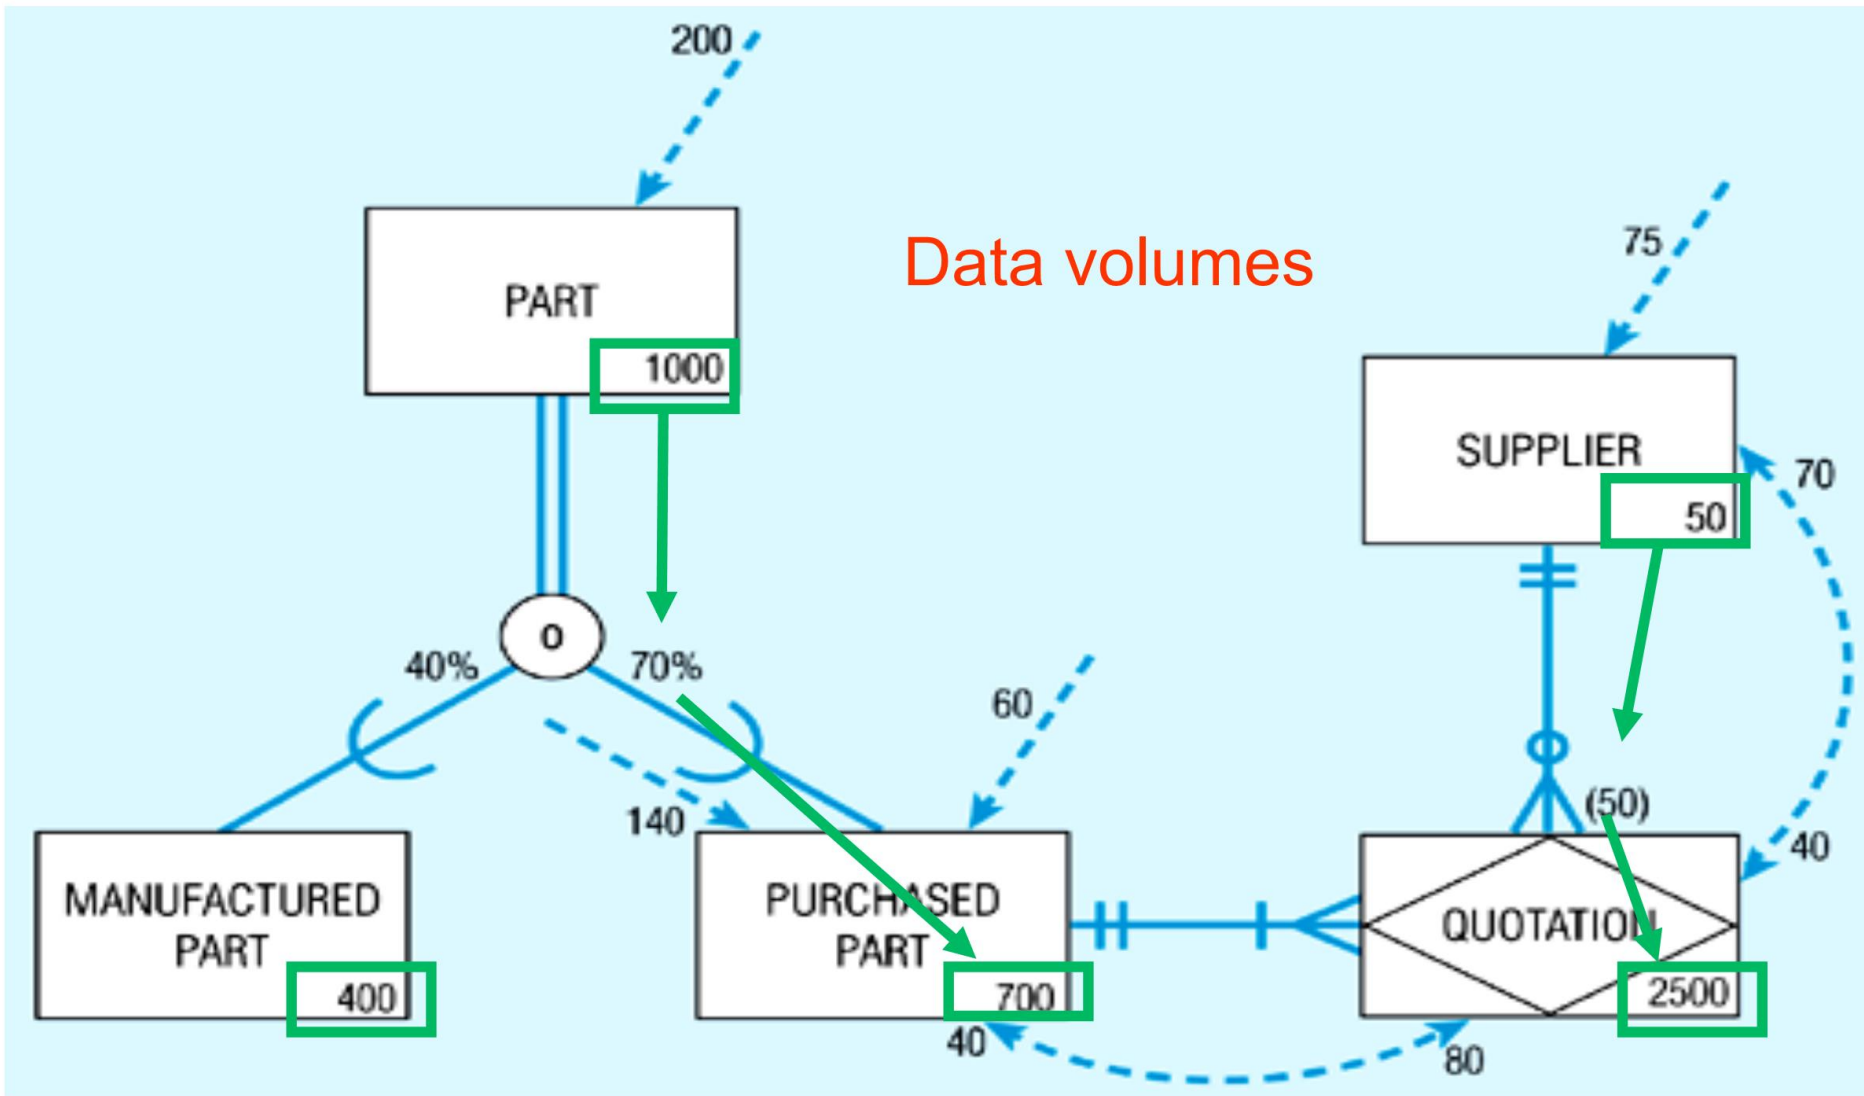


THE UNIVERSITY OF
MELBOURNE

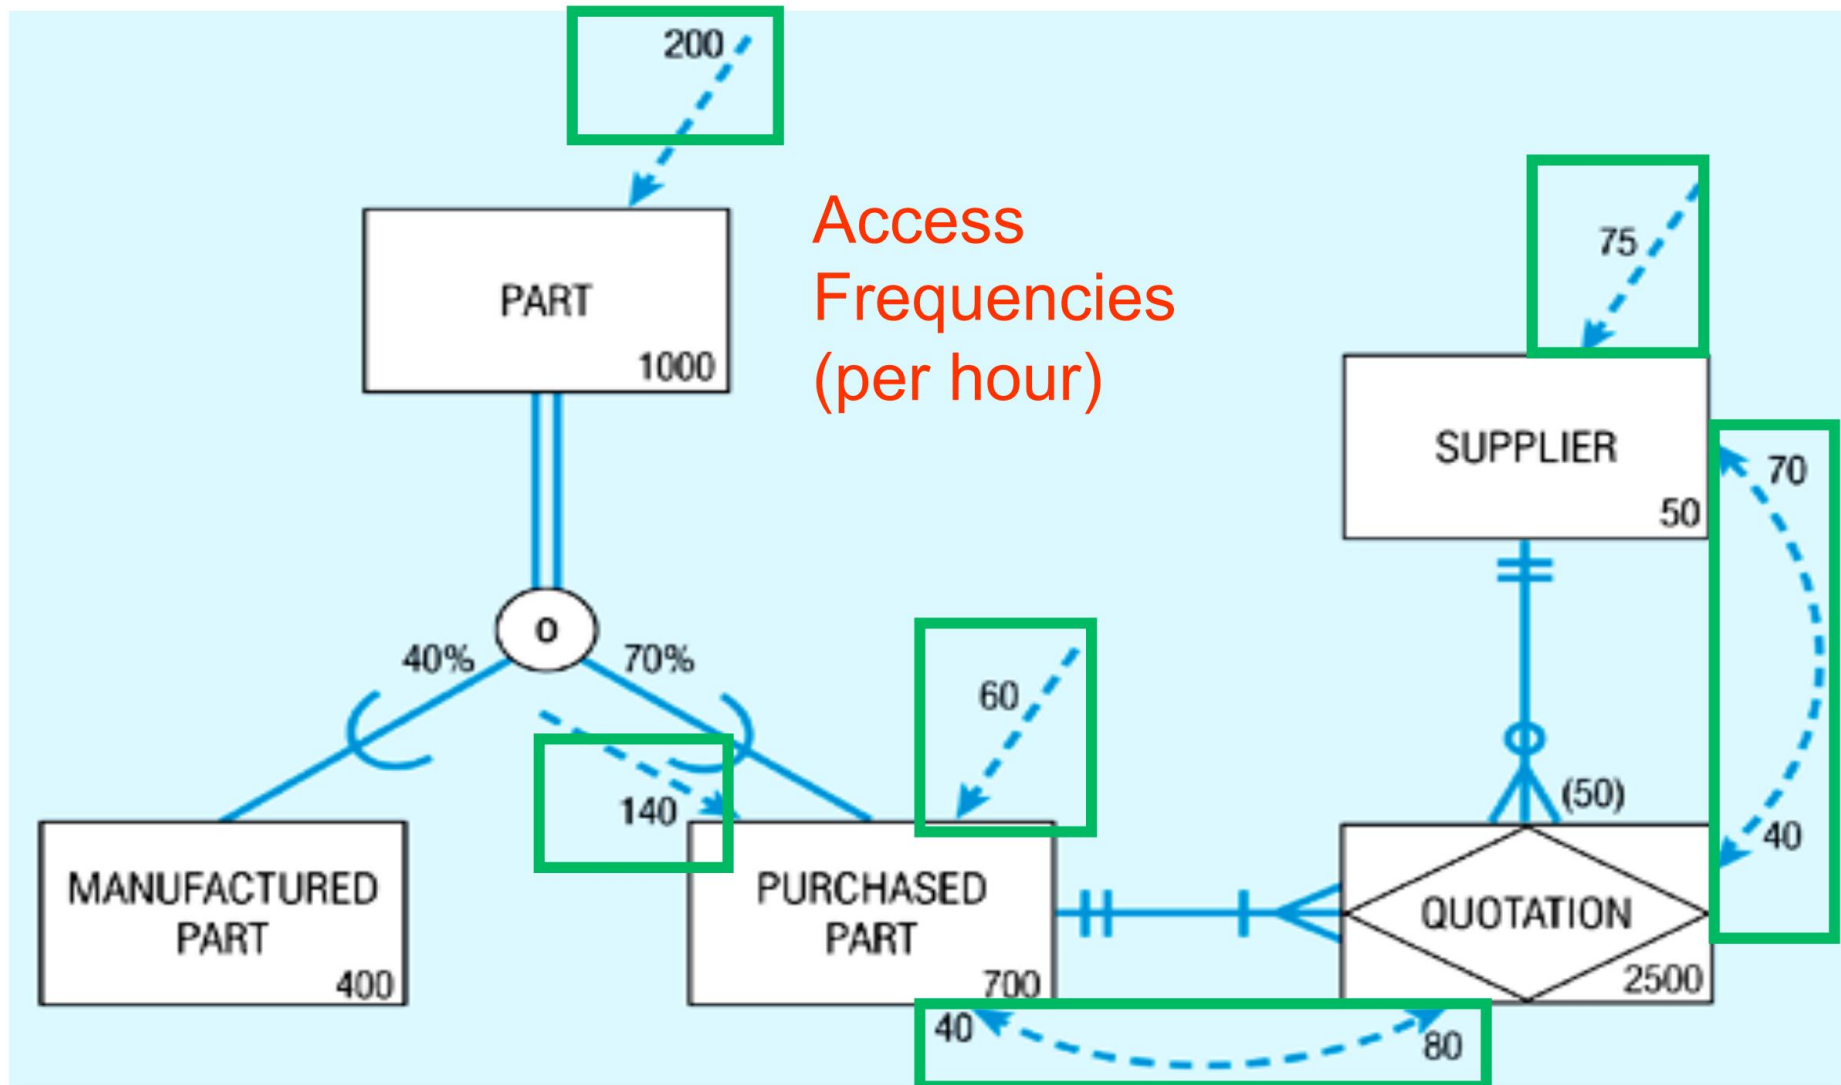
Capacity Planning

- “Capacity Planning is the process of predicting when future load levels will saturate the system and determining the most cost-effective way of delaying system saturation as much as possible.”
 - Menasce and Virgilio (2002) ‘Capacity Planning for Web Services’. Prentice Hall.
- When implementing a database, need to consider:
 - disk space requirements
 - transaction throughput
 - (at go-live and throughout the life of the system)
 - Plan for 7y Can be 20y



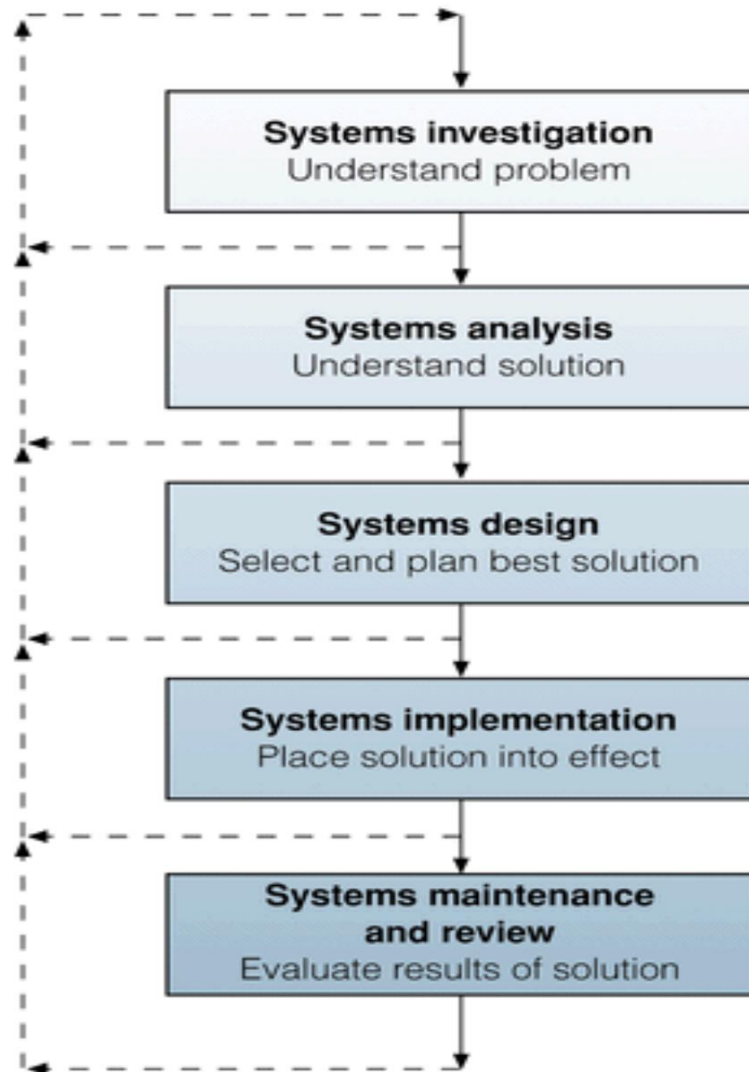


Estimating Database Usage





Capacity Planning in the dev life cycle



calculate expected
disk usage and
transaction load

monitor and predict
disk usage and
transaction load

- Which estimation methodology to use?
 - many vendors sell capacity planning solutions
 - most have the same ideas at their core
 - here we present the core concepts
- treat Database size as the sum of all Table sizes
 - where table size = number of rows * average row width

width



Id	PostedBy	Forum	Content	ParentPost	WhenPosted
1	4	NULL	April is the cruellest month, breeding	4	2015-07-23 11:00:00
2	4	3	Lilacs out of the dead land, mixing	NULL	2015-03-11 11:00:00
3	3	NULL	Memory and desire, stirring	17	2014-11-04 11:00:00
4	3	NULL	Dull roots with spring rain.	68	2015-07-29 11:00:00
5	3	NULL	Winter kept us warm, covering	38	2014-11-30 11:00:00
6	3	NULL	Earth in forgetful snow, feeding	75	2015-06-29 10:00:00
7	3	NULL	A little life with dried tubers.	6	2015-06-07 10:00:00
8	5	NULL	Summer surprised us, coming over the Starnber...	76	2015-07-20 10:00:00
9	5	NULL	With a shower of rain; we stopped in the colonn...	21	2014-12-03 11:00:00
10	4	3	And went on in sunlight, into the Hofgarten,	NULL	2015-07-21 10:00:00

height



- need to know storage size of different data types
- <https://dev.mysql.com/doc/refman/5.1/en/storage-requirements.html>

Storage Requirements for Numeric Types

Data Type	Storage Required
<u>TINYINT</u>	1 byte
<u>SMALLINT</u>	2 bytes
<u>MEDIUMINT</u>	3 bytes
<u>INT</u> , <u>INTEGER</u>	4 bytes
<u>BIGINT</u>	8 bytes
<u>FLOAT</u> (<i>p</i>)	4 bytes if $0 \leq p \leq 24$, 8 bytes if $25 \leq p \leq 53$
<u>FLOAT</u>	4 bytes
<u>DOUBLE</u> [<u>PRECISION</u>], <u>REAL</u>	8 bytes
<u>DECIMAL</u> (<i>M</i> , <i>D</i>), <u>NUMERIC</u> (<i>M</i> , <i>D</i>)	Varies; see following discussion
<u>BIT</u> (<i>M</i>)	approximately $(M+7)/8$ bytes

- <https://dev.mysql.com/doc/refman/5.1/en/storage-requirements.html>
- (these sizes are for MySQL and are slightly different for other vendors)

Storage Requirements for Date and Time Types

Data Type	Storage Required
<u>DATE</u>	3 bytes
<u>TIME</u>	3 bytes
<u>DATETIME</u>	8 bytes
<u>TIMESTAMP</u>	4 bytes
<u>YEAR</u>	1 byte

- <https://dev.mysql.com/doc/refman/5.1/en/storage-requirements.html>

Storage Requirements for String Types

In the following table, M represents the declared column length in characters for nonbinary string types and bytes for binary string types. L represents the actual length in bytes of a given string value.

Data Type	Storage Required
<code>CHAR (M)</code>	$M \times w$ bytes, $0 \leq M \leq 255$, where w is the number of bytes required for the maximum-length character in the character set. See Section 14.6.3.12.5, “Physical Row Structure” for information about <code>CHAR</code> data type storage requirements for <code>InnoDB</code> tables.
<code>BINARY (M)</code>	M bytes, $0 \leq M \leq 255$
<code>VARCHAR (M)</code> , <code>VARBINARY (M)</code>	$L + 1$ bytes if column values require 0 – 255 bytes, $L + 2$ bytes if values may require more than 255 bytes
<code>TINYBLOB</code> , <code>TINYTEXT</code>	$L + 1$ bytes, where $L < 2^8$
<code>BLOB</code> , <code>TEXT</code>	$L + 2$ bytes, where $L < 2^{16}$
<code>MEDIUMBLOB</code> , <code>MEDIUMTEXT</code>	$L + 3$ bytes, where $L < 2^{24}$
<code>LONGBLOB</code> , <code>LONGTEXT</code>	$L + 4$ bytes, where $L < 2^{32}$
<code>ENUM ('value1', 'value2', ...)</code>	1 or 2 bytes, depending on the number of enumeration values (65,535 values maximum)

- How will tables grow over time?
- Gather estimates during system analysis, e.g.
 - “The company sells 1000 products. There are 2,000,000 customers who place, on average, 5 orders each per month. An average order is for 8 different products.”

therefore:

the Product table has 1000 rows.

the Customer table has 2,000,000 rows.

the Orders table grows by 10,000,000 rows per month.

the OrderItems table grows by 80,000,000 rows per month.

- Using this simplified database as an example, assume there are:

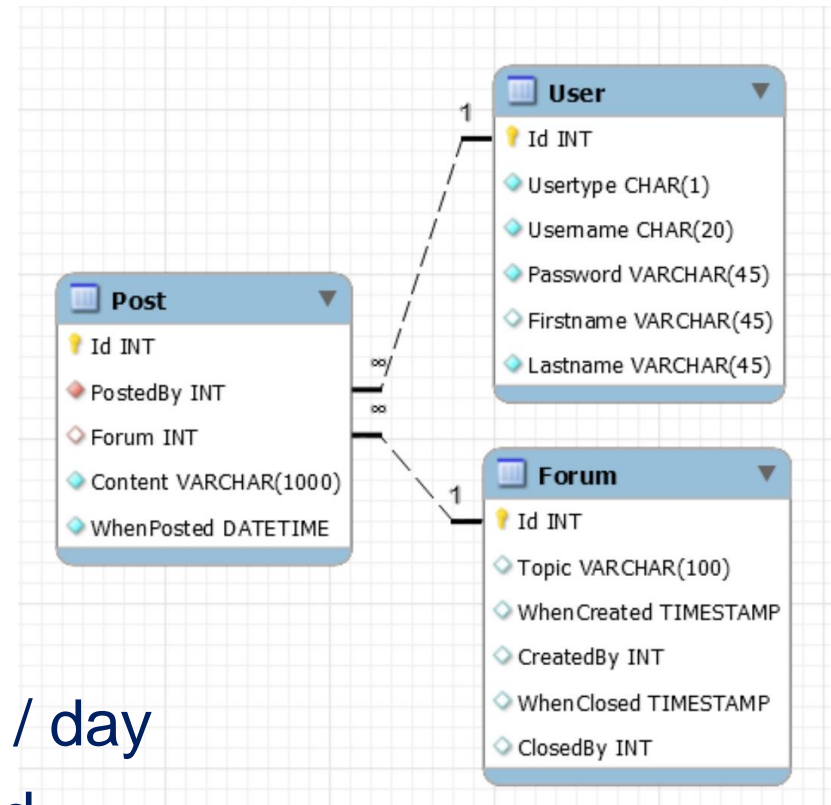
- 100 forums
- 1 million users

and assume that:

- users post average
- 30 times per month

we calculate:

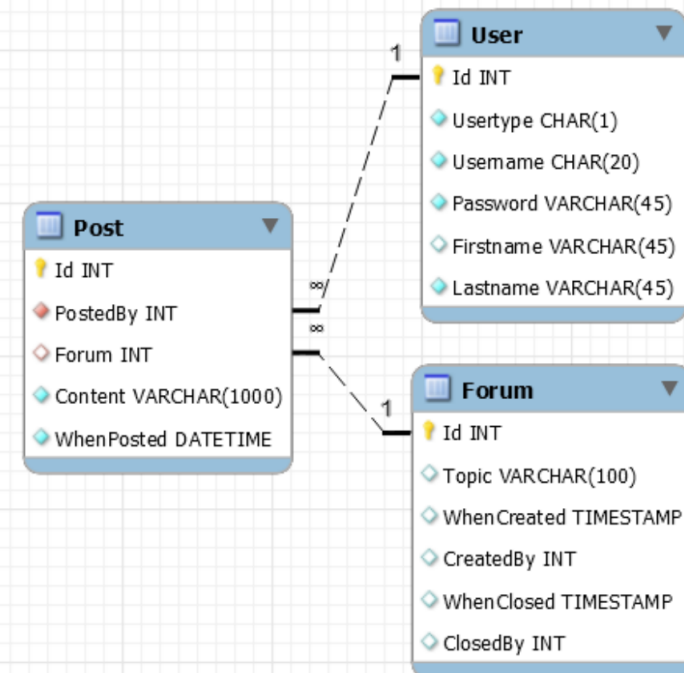
- Post table grows by 1M rows / day
- which is 12 inserts per second





Calculate disk space per table

column	type	width	rows	1 month	1 year
USER					
Id	int	4			
UserType	char(1)	1			
UserName	char(10)	10			
Password	char(10)	10			
FirstName	varchar(45)	12			
LastName	varchar(45)	15			
ROW WIDTH		52	1,000,000	1,100,000	2,000,000
DISK SPACE			52,000,000	57,200,000	104,000,000
FORUM					
Id	int	4			
Topic	varchar(100)	50		per month	
WhenCreated	timestamp	4		1	
CreatedBy	int	4			
ClosedBy	int	4			
ROW WIDTH		66	100	101	113
DISK SPACE			6,600	6,666	7,458
POST					
Id	bigint	8			
PostedBy	int	4		per user per month	
Forum	int	4		30	
Content	varchar(1000)	500			
WhenPosted	datetime	8			
ROW WIDTH		524	0	30,000,000	390,000,000
DISK SPACE			0	15,720,000,000	204,360,000,000







Projected total storage requirements

<i>Table</i>	<i>Row width</i>	<i>No. rows at 1 year</i>	<i>Size</i>
User	52 bytes	2,000,000	104 Mb
Forum	66 bytes	113	0.007 Mb
Post	524 bytes	390,000,000	204 Gb
		<i>TOTAL -></i>	<i>204 Gb</i>

Go live




1 year





- consider each business transaction
- how often will transaction each be run?
- for each transaction, what SQL statements are being run?
- for example, consider this fictitious banking application:

<i>Transaction</i>	<i>Selects</i>	<i>Inserts</i>	<i>Updates</i>	<i>Delete</i>	<i>SQL/tr</i>	<i>Tr/cust/month</i>	<i>SQL/month</i>	<i>SQL/second</i>
Withdraw	1	1	1		3	20	60,000,000	23
Deposit		1	1		2	5	10,000,000	4
Transfer	1	1	2		4	8	32,000,000	12
								39
no. customers		1,000,000						



Make an assumption
About email varchar(200)

	Go Live	Year 1	Year 2
Users	0	300000	900000
Tweets p/user	0	3600	4800
Total Tweets	0	108,000,000	432,000,000



THE UNIVERSITY OF
MELBOURNE

Backup and Recovery

- A backup is a copy of your data
 - however there are several types of backup
- If data becomes corrupted or deleted or held to ransom it can be restored from the backup copy
- A backup and recovery strategy is needed
 - To plan how data is backed up
 - To plan how it will be recovered

- human error

- e.g. accidental drop or delete
- example:

<http://www.theaustralian.com.au/australian-it/human-error-triggered-nab-software-corruption/story-e6frgakx-1225962953523>



- hardware or software malfunction

- bug in application
- hard drive (failure or corruption)
- CPU
- memory




- malicious activity
 - security compromise
 - server, database, application
- natural or man made disasters
 - consider the *scale* of the damage
- government regulation
 - historical archiving rules
 - Metadata collection (AUS)
 - HIPPA, EU data retention regulations
 - Privacy Rules

Security

Texas cops lose evidence going back eight years in ransomware attack

We have to get very, very tough on cyber and cyber warfare... and backups?

By Alexander J Martin 27 Jan 2017 at 16:57

36  SHARE ▼



🔗 I hacked the sheriff, but I did not hack his deputy 🔗

Updated Cockrell Hill, Texas has a population of just over 4,000 souls and a police force that managed to lose eight years of evidence when a departmental server was compromised by ransomware.

In a public statement, the department said the malware had been introduced to the department's systems through email. Specifically, it arrived "from a cloned email address imitating a department issued email address" and after taking root, requested 4 Bitcoin in ransom, worth about \$3,600 today, or "nearly \$4,000" as the department put it.





Failures can be divided into the following categories:

- Statement failure
 - Syntactically incorrect
- User Process failure
 - The process doing the work fails (errors, dies)
- Network failure
 - Network failure between the user and the database
- User error
 - User accidentally drops the rows, table, database
- Memory failure
 - Memory fails, becomes corrupt
- Media Failure
 - Disk failure, corruption, deletion



- Physical vs Logical
- Online vs Offline
- Full vs Incremental
- *Onsite v Offsite*

- Physical backup
 - raw copies of files and directories
 - suitable for large databases that need fast recovery
 - database is preferably offline (“cold” backup) when backup occurs
 - MySQL Enterprise automatically handles *file* locking, so database is not wholly off line
 - backup = exact copies of the database directories and files
 - backup should include logs
 - backup is only portable to machines with a similar configuration
 - to restore
 - shut down DBMS
 - copy backup over current structure on disk
 - restart DBMS

- Logical backup
 - backup completed through SQL queries
 - slower than physical
 - SQL Selects rather than OS copy
 - output is larger than physical
 - doesn't include log or config files
 - machine independent
 - server is available during the backup
 - in MySQL can use the backup using
 - Mysqldump
 - SELECT ... INTO OUTFILE
 - to restore
 - Use mysqlimport, or LOAD DATA INFILE within the mysql client



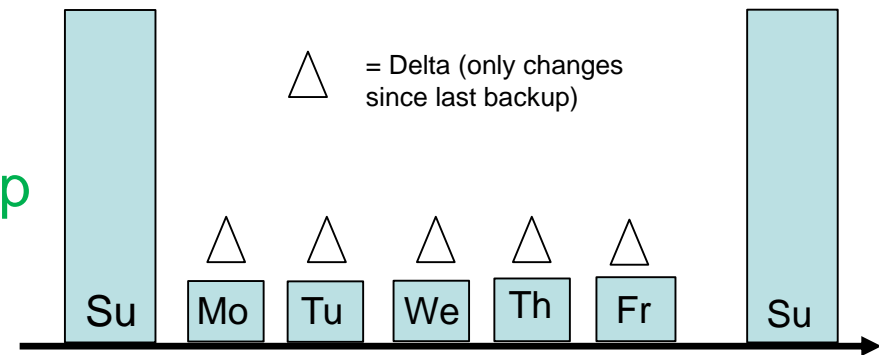
- Online (or HOT) backup
 - backups occur when the database is “live”
 - clients don’t realise a backup is in progress
 - need to have appropriate locking to ensure integrity of data
- Offline (or COLD) backup
 - backups occur when the database is stopped
 - to maximize availability to users,
take backup from replication server not live server
 - simpler to perform
 - cold backup is preferable, but not available in all situations
e.g. applications without downtime

- Full
 - a full backup is where the complete database is backed up
 - may be Physical or Logical, Online or Offline
 - it includes everything you need to get the database operational in the event of a failure
- Incremental
 - only the changes since last backup are backed up
 - for most databases this means only backup log files
 - to restore:
 - stop the database, copy backed up log files to disk
 - start the database and tell it to redo the log files

- Backup strategy is usually a combination of full and incremental backups

– for example:

- weekly full backup
- weekday incremental backup



- Conduct backups when database load is low
- If using replication, use the mirror database for backups to negate any performance concerns with the primary database
- TEST your backup before you NEED your backup!

- enables *disaster recovery*
(because backup is not physically near the disaster site)
- example solutions:
 - backup tapes transported to underground vault
 - remote mirror database maintained via replication
 - backup to Cloud (see figure below)

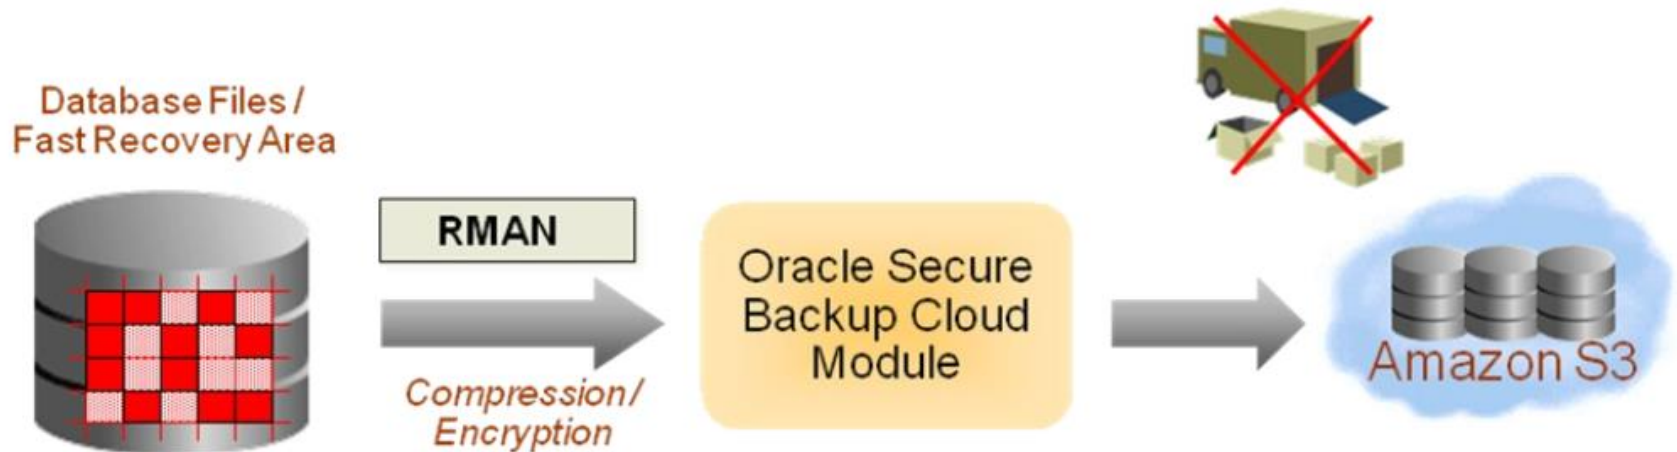


Figure 1. Oracle Database backup in the Cloud



- The roles of a DBA
 - Capacity planning
 - Calculating Capacity & Transaction workload
 - Back up and Recovery
 - Backup Types
 - Types of Failures



- Week 10:
- Lecture 19: Transactions
- Lecture 20: Data Warehousing