




# INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 07  
Relational Algebra

Semester 1 2018, Week 4



## SQL:

- Language for data manipulation
- Allow to create/delete tables, add/update/remove data, etc

Introduced next time

## Relational algebra:

- The theory behind SQL
- Makes sure that SQL produces correct answers
- Inputs/outputs are relations

Today

How do we manipulate with this data?

1. **Selection** ( $\sigma$ ): Selects a subset of *rows* from relation (horizontal filtering).
2. **Projection** ( $\pi$ ): Retains only wanted *columns* from relation (vertical filtering).
3. **Cross-product** ( $\times$ ): Allows us to combine two relations.
4. **Set-difference** ( $-$ ): Tuples in one relation, but not in the other.
5. **Union** ( $\cup$ ): Tuples in one relation and/or in the other.

Each operation returns a relation, operations can be composed

- Selection & Projection
- Union, Set Difference & Intersection
- Cross product & Joins
- Examples

*Readings: Chapter 4, Ramakrishnan & Gehrke, Database Systems*



# Example Instances

**Reserves  
(R1)**

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

**Boats**

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

**Sailors 1  
(S1)**

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

**Sailors 2  
(S2)**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



- Selection & Projection
- Union, Set Difference & Intersection
- Cross product & Joins
- Examples

*Readings: Chapter 4, Ramakrishnan & Gehrke, Database Systems*

- Retains only attributes that are in the *projection list*
- *Schema* of result:
  - Only the fields in the projection list, with the same names that they had in the input relation
- Projection operator has to *eliminate duplicates*
  - How do they arise? Why remove them?
  - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it

1. Find ages of sailors :

$\pi_{age}(S2)$

2. Find names and rating of sailors :

$\pi_{sname, rating}(S2)$

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2

sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

$\pi_{sname, rating}(S2)$

age
35.0
55.5

Removed duplicates

$\pi_{age}(S2)$



- Selects rows that satisfy *selection condition*
- Result is a relation. *Schema* of the result is same as that of the input relation.
- Do we need to do duplicate elimination?
- **Example:**

*Find sailors whose rating is above 8*

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$\sigma_{rating > 8}(S2)$



<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$\sigma_{rating > 8}(S2)$



- Conditions are standard arithmetic expressions  
 $>$ ,  $<$ ,  $\geq$ ,  $\leq$ ,  $=$ ,  $\neq$

- Conditions are combined with AND/OR clauses

And:  $\wedge$

Or:  $\vee$

- **Example:**

*Find sailors whose rating is above 8 and who are younger than 50*

$$\sigma_{rating > 8 \wedge age < 50} (S2)$$



- Operations can be combined
- Select rows that satisfy *selection condition* & retain only *certain attributes (columns)*
- **Example:**

*Find names and rating of sailors whose rating is above 8*

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



sname	rating
yuppy	9
rusty	10

$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$



- Selection & Projection
- Union, Set Difference & Intersection
- Cross product & Joins
- Examples

*Readings: Chapter 4, Ramakrishnan & Gehrke, Database Systems*



- **Union:** Combines both relations together
- **Set-difference:** Retains rows of one relation that do not appear in the other relation
- These operations take two input relations, which must be *union-compatible*:
  - Same number of fields
  - Corresponding fields have the same type

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

**S1**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

**S2**

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

**$S1 \cup S2$**

Duplicates are removed

# Set Difference

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

**S1**

sid	sname	rating	age
22	dustin	7	45.0

**S1 – S2**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

**S2**

# Set Difference

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

**S1**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

**S2**

sid	sname	rating	age
22	dustin	7	45.0

$S1 - S2$

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
44	guppy	5	35.0

$S2 - S1$

Set-difference is not symmetrical



- In addition to the 5 basic operators, there are several additional “**Compound Operators**”
  - These add no computational power to the language, but are useful shorthands
  - Can be expressed solely with the basic operations
- **Intersection** retains rows that appear in *both* relations
- Intersection takes two input relations, which must be *union-compatible*.
- Q: How to express it using basic operators?

$$R \cap S = R - (R - S)$$



## Example:

*Find sailors who appear in both relations S1 and S2*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

**S1**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

**S2**

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

**$S1 \cap S2$**



- Selection & Projection
- Union, Set Difference & Intersection
- Cross product & Joins
- Examples

*Readings: Chapter 4, Ramakrishnan & Gehrke, Database Systems*

- **Cross product** combines two relations:
  - Each row of one input is merged with each row from another input
  - Output is a new relation with all attributes of *both* inputs
  - $\times$  is used to denote cross-product
- Example:  $S1 \times R1$ 
  - Each row of  $S1$  paired with each row of  $R1$
- Question: How many rows are in the result?
  - A:  $\text{card}(S1) * \text{card}(R1)$



# Cross Product Example

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

**S1**

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

**R1**

**S1 X R1 =**

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

- *Result schema* has one field per field of S1 and R1, with field names “inherited” if possible.
  - May have a *naming conflict*, i.e. both S1 and R1 have a field with the same name (e.g. *sid*).
  - In this case, can use the *renaming operator*.

$$\rho(C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$$

Result relation

C	sid1	sname	rating	age	sid2	bid	day
	22	dustin	7	45.0	22	101	10/10/96
	22	dustin	7	45.0	58	103	11/12/96
	31	lubber	8	55.5	22	101	10/10/96
	31	lubber	8	55.5	58	103	11/12/96
	58	rusty	10	35.0	22	101	10/10/96
	58	rusty	10	35.0	58	103	11/12/96

- Joins are compound operators involving cross product, selection, and (sometimes) projection.
- Most common type of join is a **natural join** (often just called **join**).  $R \bowtie S$  conceptually is a cross product that matches rows where attributes that appear in both relations have equal values (and we omit duplicate attributes).
- To obtain cross product a DBMS must:
  1. Compute  $R \times S$
  2. Select rows where attributes that appear in both relations have equal values
  3. Project all unique attributes and one copy of each of the common ones.



## Example:

*Find all sailors (from relation S1) who have reserved a boat*

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

R1

$S1 \bowtie R1 =$

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96



# Natural Join Example

1

**S1 X R1 =**

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

# Natural Join Example

1

**S1 X R1 =**

2

**σ**

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
<del>22</del>	<del>dustin</del>	<del>7</del>	<del>45.0</del>	<del>58</del>	<del>103</del>	<del>11/12/96</del>
<del>31</del>	<del>lubber</del>	<del>8</del>	<del>55.5</del>	<del>22</del>	<del>101</del>	<del>10/10/96</del>
<del>31</del>	<del>lubber</del>	<del>8</del>	<del>55.5</del>	<del>58</del>	<del>103</del>	<del>11/12/96</del>
<del>58</del>	<del>rusty</del>	<del>10</del>	<del>35.0</del>	<del>22</del>	<del>101</del>	<del>10/10/96</del>
58	rusty	10	35.0	58	103	11/12/96

# Natural Join Example

①  $S1 \times R1 =$

②  $\sigma$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
<del>22</del>	<del>dustin</del>	<del>7</del>	<del>45.0</del>	<del>58</del>	<del>103</del>	<del>11/12/96</del>
<del>31</del>	<del>lubber</del>	<del>8</del>	<del>55.5</del>	<del>22</del>	<del>101</del>	<del>10/10/96</del>
<del>31</del>	<del>lubber</del>	<del>8</del>	<del>55.5</del>	<del>58</del>	<del>103</del>	<del>11/12/96</del>
<del>58</del>	<del>rusty</del>	<del>10</del>	<del>35.0</del>	<del>22</del>	<del>101</del>	<del>10/10/96</del>
58	rusty	10	35.0	58	103	11/12/96

$\pi$  ③

$S1 \bowtie R1 =$

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

- **Condition Join (or theta-join)** is a cross product with a condition.

$$R \bowtie_c S = \sigma_c (R \times S)$$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

–Result schema is the same as that of cross-product

- **Equi-Join** is a special case of condition join, where condition  $c$  contains only *equalities* (e.g.  $S1.sid = R1.sid$ )
  - Is this then a natural join? What is different?



- Selection & Projection
- Union, Set Difference & Intersection
- Cross product & Joins
- Examples

*Readings: Chapter 4, Ramakrishnan & Gehrke, Database Systems*



## Boats

<u>bid</u>	bname	color
101	Interlake	Blue
102	Interlake	Red
103	Clipper	Green
104	Marine	Red

## Sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

## Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

*Find names of sailors who have reserved boat #103*

Solution 1:  $\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$

Solution 2:  $\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$

*Find names of sailors who have reserved a blue boat*

- Information about boat color only available in Boats; so need an extra join:

$$\pi_{sname}((\sigma_{color='blue'}Boats) \bowtie Reserves \bowtie Sailors)$$

- A more efficient solution:

$$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color='blue'}Boats) \bowtie Res.) \bowtie Sailors)$$

You shouldn't worry about efficiency for now. A DBMS will do this job for us and find alternative 2.





Find all pairs of sailors in which the older sailor has a lower rating



1. Find (the name of) all sailors whose rating is above 9
2. Find all sailors who reserved a boat prior to November 1, 1996
3. Find (the names of) all boats that have been reserved at least once
4. Find all pairs of sailors with the same rating



- Relational Algebra Operations: Selection, Projection, Union, Set, Difference, Intersection, **JOINS**...
- Draw different queries with Relational Algebra operations



- Introducing SQL