



INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 9
SQL Summary

Semester 1 2018, Week 5



1. Find the name of all sailors whose rating is above 9

$$\rho_{sname}(S_{rating>9}(Sailors))$$

2. Find all sailors who reserved a boat prior to November 1, 1996

$$\rho_{sname}(Sailors \bowtie S_{day<'11/1/96'}(Reserves))$$

3. Find (the names of) all boats that have been reserved at least once

$$\rho_{bname}(Boats \bowtie Reserves)$$



4. Find all pairs of sailors with the same rating

$r(S1(1 \rightarrow sid1, 2 \rightarrow sname1, 3 \rightarrow rating1, 4 \rightarrow age1), Sailors)$

$r(S2(1 \rightarrow sid2, 2 \rightarrow sname2, 3 \rightarrow rating2, 4 \rightarrow age2), Sailors)$

$\rho_{sname1, sname2}(S1 \bowtie_{rating1=rating2} S2)$



- Do I need to make assumptions about **all** possible extensions to the requirements in the future?
 - No. Your model should capture the existing description of the problem
- What makes the best solution?
 - A solution that is valid (captures the existing set of requirements) but is *flexible* that for each requirement will allow some *minor* extensions.
 - E.g. Instead of 3 reaction gravities (strong, moderate, weak), you get an additional one
- Do I need to use supertype/subtype to score the highest?
 - No. If your model is valid and flexible – you will get the highest marks in both cases (with or without hierarchies)

Examples for marks deduction:

- Entity / Attribute incorrect or missing
- Relationship cardinality incorrect (e.g. one - many)
- Poor naming of object (e.g. T1, T2)
- Wrong data type (e.g. varchar for placing 'age')
- Incorrect primary or foreign key
- Not Null is wrong
- Unresolved M-M exists, or associative entity is incorrect
- Business rules can't be supported (e.g. one medicare card cannot capture the entire family)

NOTE: These are EXAMPLES, not a contract set in stone



- Extending your knowledge
 - DML
 - Comparison & Logic Operators
 - Set Operations
 - Subquery
 - Multiple record INSERTs
 - INSERT from a table, UPDATE, DELETE, REPLACE
 - Views
 - DDL
 - ALTER and DROP, TRUNCATE, RENAME
- How to think about SQL
 - Problem Solving



- SQL keywords are case insensitive
 - We try to CAPITALISE them to make them clear
- Table names are Operating System Sensitive
 - If case sensitivity exists in the operating system, then the table names are case sensitive! (i.e. Mac, Linux)
 - Account <> ACCOUNT
- Field names are case insensitive
 - ACCOUNTID == AccountID == AcCoUnTID
- You can do maths in SQL...
 - SELECT 1*1+1/1-1;



- Comparison:


Operator	Description
=	Equal to
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to
<> OR !=	Not equal to (depends on DBMS as to which is used)

- Logic:


- AND, NOT, OR
- **Example:** SELECT * FROM Furniture WHERE ((Type="Chair" AND Colour = "Black") OR (Type = "Lamp" AND Colour = "Black"))




- UNION
 - Shows all rows returned from the queries (or tables)
- INTERSECT
 - Shows only rows that are common in the queries (or the tables)
- [UNION/INTERSECT] ALL
 - If you want duplicate rows shown in the results you need to use the ALL keyword.. UNION ALL etc.
- In MySQL only UNION and UNION ALL are supported




- T1 INNER JOIN T2 ON T1.ID = T2.ID
- T1 NATURAL JOIN T2




- T1 LEFT JOIN T2 ON T1.ID = T2.ID



- T1 RIGHT JOIN T2 ON T1.ID = T2.ID



- T1 **FULL OUTER JOIN** T2 **ON** T1.ID = T2.ID

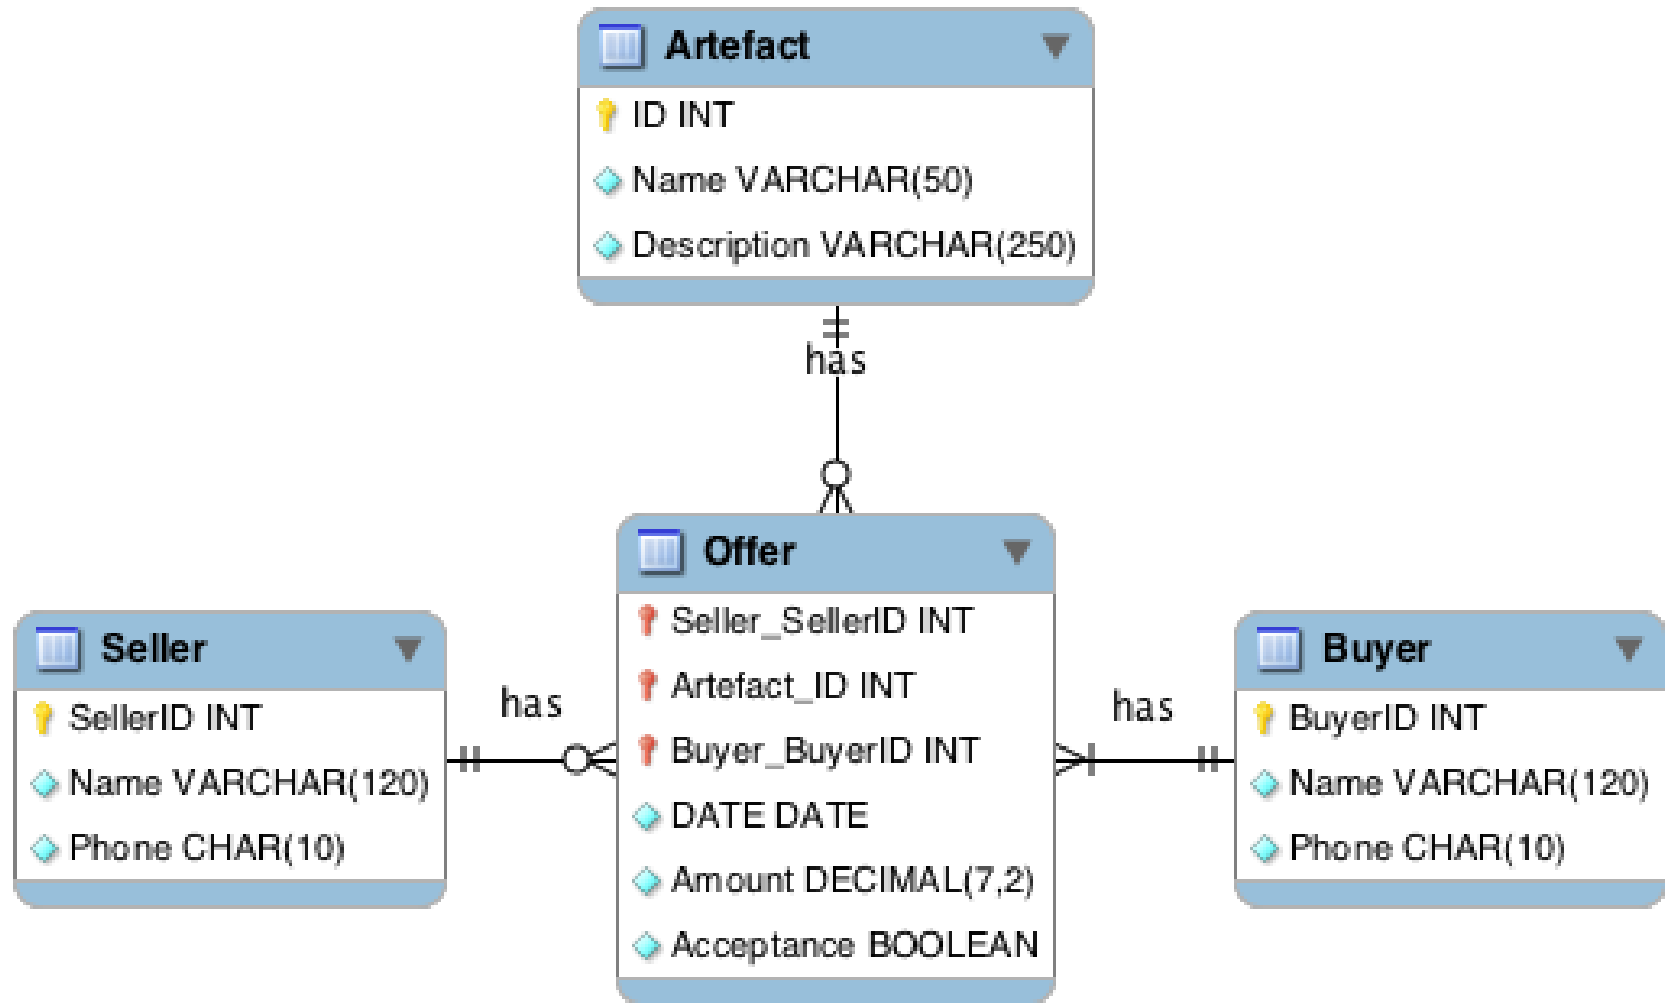




- SQL provides the ability to *nest* subqueries
- A nested query is simply another select query you write to produce a table set
 - Remember that all select queries return a table set of data
- A common use of subqueries is to perform set tests
 - Set membership, set comparisons



- IN / NOT IN
 - Used to test whether the attribute is IN/NOT IN the subquery list
- ANY
 - True if any value returned meets the condition
- ALL
 - True if all values returned meet the condition
- For more info: https://www.w3schools.com/sql/sql_any_all.asp





Seller

SellerID	Name	Phone
1	Abby	0233232232
2	Ben	0311111111
3	Carl	0333333333

Offer

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N

Artefact

ID	Name	Description
1	Vase	Old Vase
2	Knife	Old Knife
3	Pot	Old Pot

Buyer

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555



- List the BuyerID, Name and Phone number for all bidders on artefact 1

```
SELECT * FROM Buyer
  WHERE BuyerID IN
    (SELECT BuyerID FROM Offer WHERE ArtefactID = 1)
```

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444

- Which Artefacts don't have offers made on them

```
SELECT * FROM Artefact
WHERE ID NOT IN
(SELECT ArtefactID FROM Offer);
```

ID	Name	Description
3	Pot	Old Pot

- Which Buyers haven't made a bid for Artefact 3

```
SELECT * FROM Buyer
WHERE BuyerID NOT IN
(SELECT BuyerID FROM Offer
WHERE ArtefactID = 3);
```

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555



Which Buyers haven't made a bid for the "Pot" Artefact ?




- List the BuyerID, Name and Phone number for all bidders on artefact 1

```
SELECT * FROM Buyer
WHERE BuyerID IN (SELECT BuyerID FROM Offer
                  WHERE ArtefactID = 1)
```

Equals to

```
SELECT BuyerID, Name and Phone
FROM Buyer NATURAL JOIN Offer
WHERE ArtefactID = 1
```

This is a more efficient way





- Inserting records from a table:
 - Note: table must already exist

```
INSERT INTO NewEmployee  
SELECT * FROM Employee;
```

- Multiple record inserts:

All columns must be inserted

```
INSERT INTO Employee VALUES  
(DEFAULT, "A", "A's Addr", "2012-02-02", NULL, "S"),  
(DEFAULT, "B", "B's Addr", "2012-02-02", NULL, "S"),  
(DEFAULT, "C", "C's Addr", "2012-02-02", NULL, "S");
```

Specific columns will be inserted

```
INSERT INTO Employee  
(Name, Address, DateHired, EmployeeType)  
VALUES  
("D", "D's Addr", "2012-02-02", "C"),  
("E", "E's Addr", "2012-02-02", "C"),  
("F", "F's Addr", "2012-02-02", "C");
```


- Changes *existing* data in tables
 - Order of statements is important
 - Specifying a WHERE clause is important
 - Unless you want it to operate on the whole table

```
UPDATE Hourly  
    SET HourlyRate = HourlyRate * 1.10;
```

- **Example:** *Increase all salaries greater than \$100000 by 10% and all other salaries by 5%*

```
UPDATE Salaried  
    SET AnnualSalary = AnnualSalary * 1.05  
    WHERE AnnualSalary <= 100000;  
UPDATE Salaried  
    SET AnnualSalary = AnnualSalary * 1.10  
    WHERE AnnualSalary > 100000;
```

Any problems with this?



The UPDATE Statement: CASE

- A better solution in this case is to use the **CASE** command

```
UPDATE Salaried
SET AnnualSalary =
CASE
    WHEN AnnualSalary <= 100000
    THEN AnnualSalary * 1.05
    ELSE AnnualSalary * 1.10
END;
```

If salary is lower than 100000 increase it by 5%,
otherwise increase it by 10%



- REPLACE
 - REPLACE works identically as INSERT
 - Except if an old row in a table has a key value the same as the new row then it is overwritten...
- DELETE
 - The DANGEROUS command – deletes *ALL* records

```
DELETE FROM Employee;
```

 - The better version (unless you are really, really sure)

```
DELETE FROM Employee  
WHERE Name = "Grace";
```
 - Be aware of the foreign key constraints
 - ON DELETE CASCADE or ON DELETE RESTRICT (lab practice)

- Any relation that is not in the physical models, but is made available to the “user” as a virtual relation is called a view.
- Views are good because:
 - They help hide the query complexity from users
 - They help hide data from users
 - Different users use different views
 - Prevents someone from accessing the employee tables to see salaries for instance
 - One way of improving database security
- Create view statement:
CREATE VIEW nameofview **AS** validsqlstatement
- Once a view is defined
 - Its definition is stored in the database (not the data, but metadata – schema information)
 - Can be used just like any other table



```
CREATE VIEW EmpPay AS
SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, HourlyRate AS Pay
FROM Employee INNER JOIN Hourly
ON Employee.ID = Hourly.ID

UNION

SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, AnnualSalary AS Pay
FROM Employee INNER JOIN Salaried
ON Employee.ID = Salaried.ID

UNION

SELECT Employee.ID, Employee.Name, DateHired,
       EmployeeType, BillingRate AS Pay
FROM Employee INNER JOIN Consultant
ON Employee.ID = Consultant.ID;
```




SELECT * FROM EmpPay;

ID	Name	DateHired	EmployeeType	Pay
3	Alice	2012-12-02	H	23.43
4	Alan	2010-01-22	H	29.43
1	Sean	2012-02-02	S	92000.00
2	Linda	2011-06-12	S	92300.00
5	Peter	2010-09-07	C	210.00
6	Rich	2012-05-19	C	420.00

**SELECT * FROM EmpPay
WHERE EmployeeType = "H" OR EmployeeType = "C"**

ID	Name	DateHired	EmployeeType	Pay
3	Alice	2012-12-02	H	23.43
4	Alan	2010-01-22	H	29.43
5	Peter	2010-09-07	C	210.00
6	Rich	2012-05-19	C	420.00

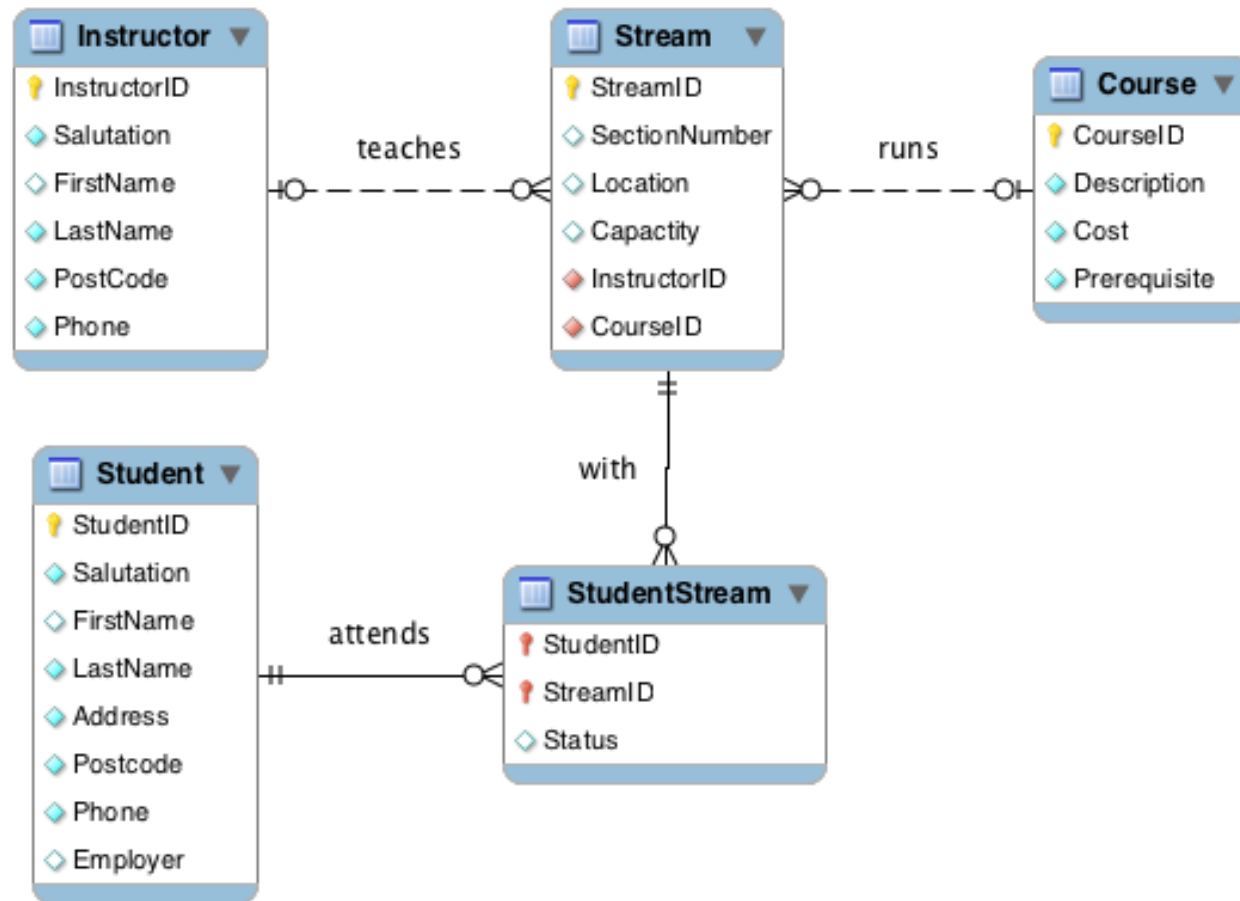
Output Snippets Query 1 Result Lecture7.sql Result x

Fetch 4 records. Duration: 0.000 sec, fetched 4 records.

- There are more than CREATE!
- ALTER
 - Allows us to add or remove attributes (columns) from a relation (table)
 - **ALTER TABLE** TableName **ADD** AttributeName AttributeType
 - **ALTER TABLE** TableName **DROP** AttributeName
- RENAME
 - Allows the renaming of tables (relations)
 - **RENAME TABLE** CurrentTableName **TO** NewTableName



- It's going to be critical for you to think like SQL to handle the queries you will need to write...
- Hopefully the following discussion will help you in this endeavour:
 1. **USE** the database design as a MAP to help you when you are formulating queries
 2. **USE** the structure of the **SELECT** statement as a template
 3. **FILL** out parts of the **SELECT** structure and **BUILD** the query
- Let's try it!

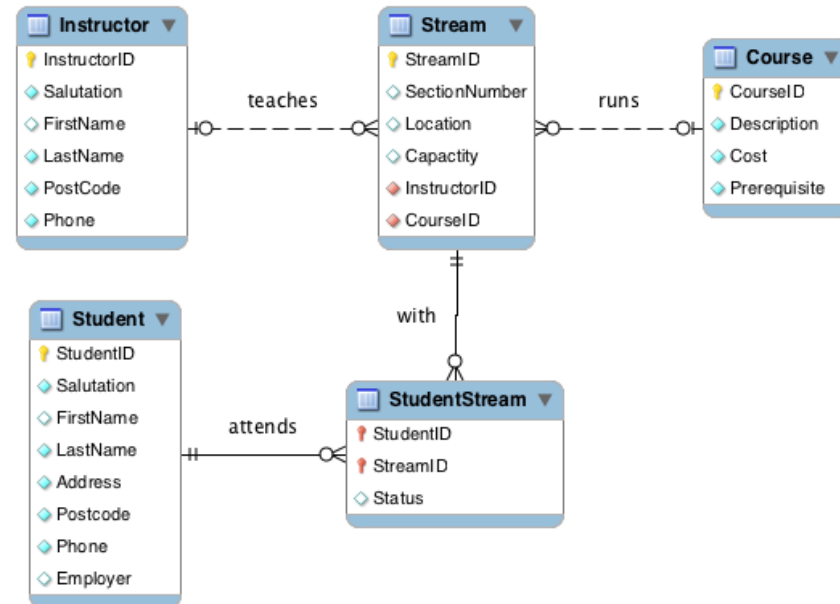


Example: Which employers employ students who are doing a course in locations where the capacity is greater than 20 persons, and what are those locations?



*Which employers employ students
who are doing a course in locations
where the capacity is greater than 20 persons,
and what are those locations?*

- What is the query asking for:
 - Which fields & tables:
F: Employer, Location, Capacity
T: Student, Stream, StudentStream
 - But only if the capacity > 20 (condition)



- Lets try to use the structure of the SELECT statement now:

```

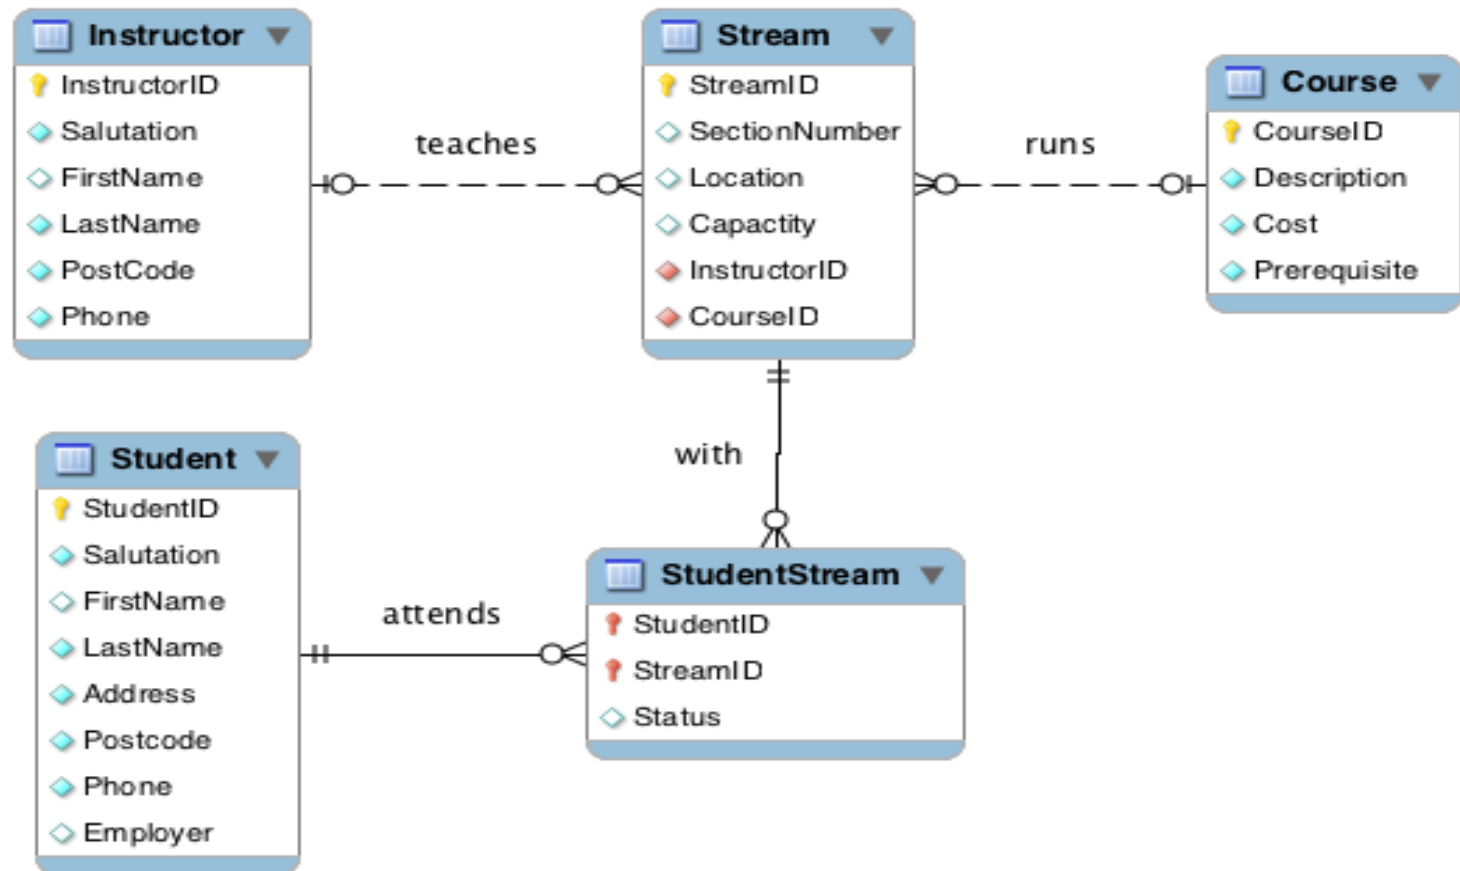
SELECT Employer, Location, Capacity
FROM Student INNER JOIN StudentStream
ON Student.StudentID = StudentStream.StudentID
INNER JOIN Stream
ON StudentStream.StreamID = Stream.StreamID
WHERE Capacity > 20;
    
```

```

SELECT Employer, Location, Capacity
FROM Student
NATURAL JOIN StudentStream
NATURAL JOIN Stream
WHERE Capacity > 20;
    
```



What is the phone number of the instructor who teaches a course that costs over 10000\$ attended by studentID 202.





- You need to know how to write SQL



- Storage and indexing
 - Learn how data is stored and accessed within a DBMS
 - Alternative types of indexes
 - Going “under the hood” of a DBMS