



INFO20003 Database Systems


Dr Renata Borovica-Gajic

Lecture 12
Query Processing Part II

Semester 1 2018, Week 6

Remember this? Components of a DBMS

Will briefly
touch upon ...



**TODAY
Joins**

This is one of several possible architectures;
each system has its own slight variations.



- Nested loops join
- Sort-merge join
- Hash join
- General joins

Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems

- Are very common and can be **very** expensive (cross product in the worst case)
- There are many implementation techniques for join operations
- Join techniques we will cover:
 1. Nested-loops join
 2. Sort-merge join
 3. Hash join

Example: SELECT *
FROM Reserves R1, Sailors S1
WHERE R1.sid=S1.sid

- In algebra: $R \bowtie S$. They are very common and need to be carefully optimized.
- $R \times S$ is large; so, $R \times S$ followed by a selection is inefficient.

Left / Outer

R	
11	80
13	75
12	10
15	80
15	44
13	74

Right / Inner

S	
11	20
13	20
12	20
13	75
13	35
12	10

- Join is associative and commutative:
 - $A \times B == B \times A$
 - $A \times (B \times C) == (A \times B) \times C$
- **Cost metric** : Number of pages; Number of I/O

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)
Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- **Sailors (S):**

- 80 tuples per page, **500 pages**
- $NPages(S) = 500$, $NTuplesPerPage(S) = 80$
- $NTuples(S) = 500 * 80 = 40000$


- **Reserves (R):**

- 100 tuples per page, **1000 pages**
- $NPages(R) = 1000$, $NTuplesPerPage(R) = 100$
- $NTuples(R) = 100000$

- For each tuple in the *outer* relation R, we scan the entire *inner* relation S

Pseudo code:

```
foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```



- Cost:

$$\text{Cost (SNJL)} = \text{NPages(Outer)} + \text{NTuples(Outer)} * \text{NPages(Inner)}$$

- Our example:

$$\begin{aligned} \text{Cost (SNLJ)} &= 1000 + 100 * 1000 * 500 \\ &= 50001000 \text{ (I/O)} \end{aligned}$$




- For each **page** of R
 - get each **page** of S
 - write out matching pairs of tuples $\langle r, s \rangle$, where r is in R-page and S is in S-page

Pseudo code:

```

foreach page  $b_R$  in R do
  foreach page  $b_S$  in S do
    foreach tuple  $r$  in  $b_R$  do
      foreach tuple  $s$  in  $b_S$  do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
    
```




$$\text{Cost (PNJL)} = \text{NPages(Outer)} + \text{NPages(Outer)} * \text{NPages(Inner)}$$

- Our example:

$$\text{Cost (PNLJ)} = 1000 + 1000 * 500 = 501000 \text{ (I/O)}$$

- Page-oriented NL doesn't exploit extra memory buffers
- **Alternative approach:**
 - Use one page as an input buffer for scanning the inner S, one page as the output buffer, and use all remaining pages to hold 'block' of outer R
- For each matching tuple r in R-block, s in S-page, add $\langle r, s \rangle$ to result. Then read next R-block, scan S, etc



$$\text{Cost (BNJL)} = \text{NPages(Outer)} + \text{NBlocks(Outer)} * \text{NPages(Inner)}$$

$$\bullet \text{ NBlocks(Outer)} = \left\lceil \frac{\text{NPages(Outer)}}{\text{Blocksize}-2} \right\rceil$$

	R	S
B1	11,80	11,20
	13,75	13,20
	12,10	12,20
B2	15,80	13,75
	15,44	13,35
	13,74	12,10

• Our example:

Let's say we have 102 pages of space in memory, and consider Reserves (R) as the outer and Sailors (S) as the inner table.

$$\text{NBlocks(R)} = 1000 / (102 - 2) = 10$$

$$\text{Cost(BNLJ)} = 1000 + 10 * 500 = 6000 \text{ I/O}$$

- Nested loops join
- Sort-merge join
- Hash join
- General joins

Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems

Sort-Merge Join ($R \bowtie S$)
 $i=j$

- **Sort** R and S on the join column, then scan them to do a **merge** (on join column), and output result tuples

R	S
11,80	11,20
12,10	12,20
13,74	12,10
13,75	13,75
15,44	13,35
15,80	13,20

- Sorted R is scanned once;
- Each S group of the same key values is scanned once per matching R tuple (typically means Sorted S is scanned once too).
- Useful when:
 - one or both inputs are already sorted on join attribute(s)
 - output is required to be sorted on join attributes(s)

$$\text{Cost (SMJ)} = \text{Sort(Outer)} + \text{Sort(Inner)} \\ + \text{NPages(Outer)} + \text{NPages(Inner)}$$

Sort inputs
Merge inputs

$$\text{Sort(R)} = \text{External Sort Cost} = 2 * \text{NumPasses} * \text{NPages(R)}$$

Our example:

Let's say that both Reserves and Sailors can be sorted in 2 passes, then:


$$\begin{aligned} \text{Cost(SMJ)} &= \text{Sort R} + \text{Sort S} + \text{NPages(R)} + \text{NPages(S)} \\ &= 2 * 2 * \text{NPages(R)} + 2 * 2 * \text{NPages(S)} \\ &\quad + \text{NPages(R)} + \text{NPages(S)} \\ &= 5 * 1000 + 5 * 500 = 7500 \text{ I/O} \end{aligned}$$

- Nested loops join
- Sort-merge join
- Hash join
- General joins

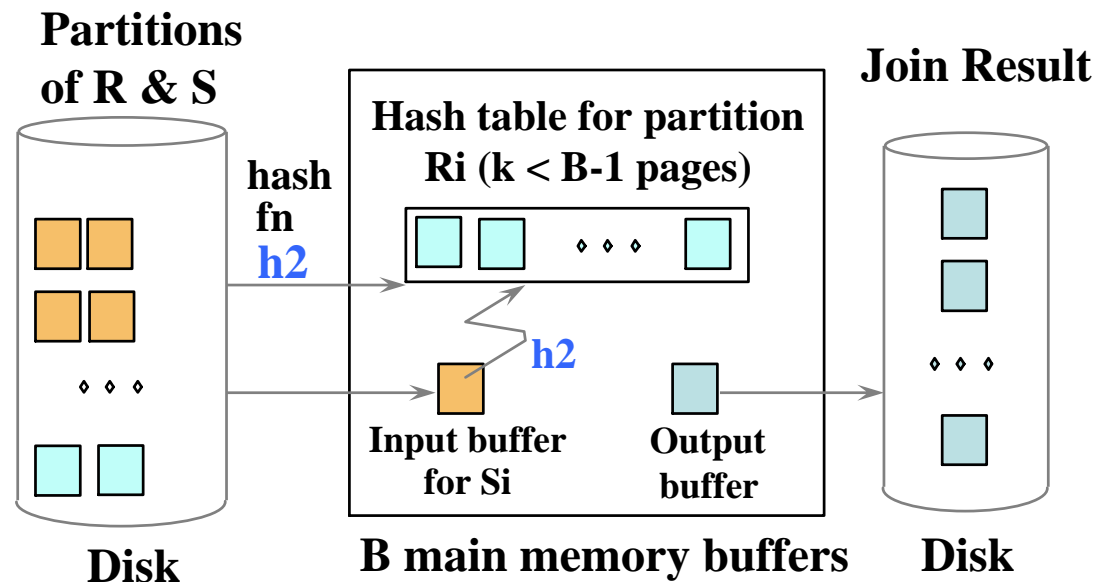
Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems

Hash-Join

- Partition both relations using hash function h :
R tuples in partition I will **only** match S tuples in partition I



- Read in a partition of R, hash it using h_2 ($\neq h$). Scan matching partition of S, probe hash table for matches





1. In partitioning phase, we read+write both relations
2. In matching phase, we read both relations

$$\begin{aligned} \text{Cost (HJ)} = & 2 * \text{NPages(Outer)} + 2 * \text{NPages(Inner)} && \text{Create partitions} \\ & + \text{NPages(Outer)} + \text{NPages(Inner)} && \text{Match partitions} \end{aligned}$$

- **Our example:**

$$\begin{aligned} \text{Cost(HJ)} &= 2 * \text{NPages(R)} + 2 * \text{NPages(S)} + \text{NPages(R)} + \text{NPages(S)} \\ &= 3 * 1000 + 3 * 500 = 4500 \text{ I/Os} \end{aligned}$$



<https://www.youtube.com/watch?v=o1dMJ6-CKzU>

From 0:58



- Relation A: 10 tuples per page, 100 pages
- Relation B: 20 tuples per page, 500 pages
- Memory Buffer 52 pages, imagine SMJ can be done in 3 passes<- not true but suffices for the purpose of testing the formula
- Find the cheapest join? Assume A is always outer, B is inner.
- Answer?

- Nested loops join
- Sort-merge join
- Hash join
- General joins

Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems

- Equalities over several attributes (e.g., *R.sid=S.sid AND R.rname=S.sname*):
 - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns
- Inequality conditions (e.g., *R.rname < S.sname*):
 - Hash Join, Sort Merge Join not applicable
 - Block NL quite likely to be the best join method here

- A virtue of relational DBMSs:
 - Queries are composed of a few basic operators
 - Implementation of operators can be **carefully tuned**
 - Important to do this
- Many alternative implementations for each operator
 - No universally superior technique for most operators
- Must consider alternatives for each operation in a query and choose best one based on system statistics...
 - Part of the broader task of optimizing a query composed of several operations



- Understand alternatives for join operator implementations
 - Be able to calculate the cost of alternatives
- Important for Assignment 3 as well

Faculty of Science Panel Event

Are you a student in the Bachelor of Science? Do you want to know where your degree can take you and how to get there?

On **Tuesday 17th April** a panel of experts will discuss how to best prepare for the transition from university into the workplace and cover topics such as:

- How to get involved
- Skill building in curriculum and extra-curricular development
- Recognise and reflect on your development
- Networking options and peer reflection
- Necessary skills in a STEM workplace

Tuesday 17 April, 1:15pm – 2:15pm
Harold White Theatre, 757 Swanston St

Registration here:

<https://form.jotform.com/80978604502965>





- Query optimization
 - How does a DBMS pick a good query plan?