



INFO20003 Database Systems

Dr Renata Borovica-Gajic

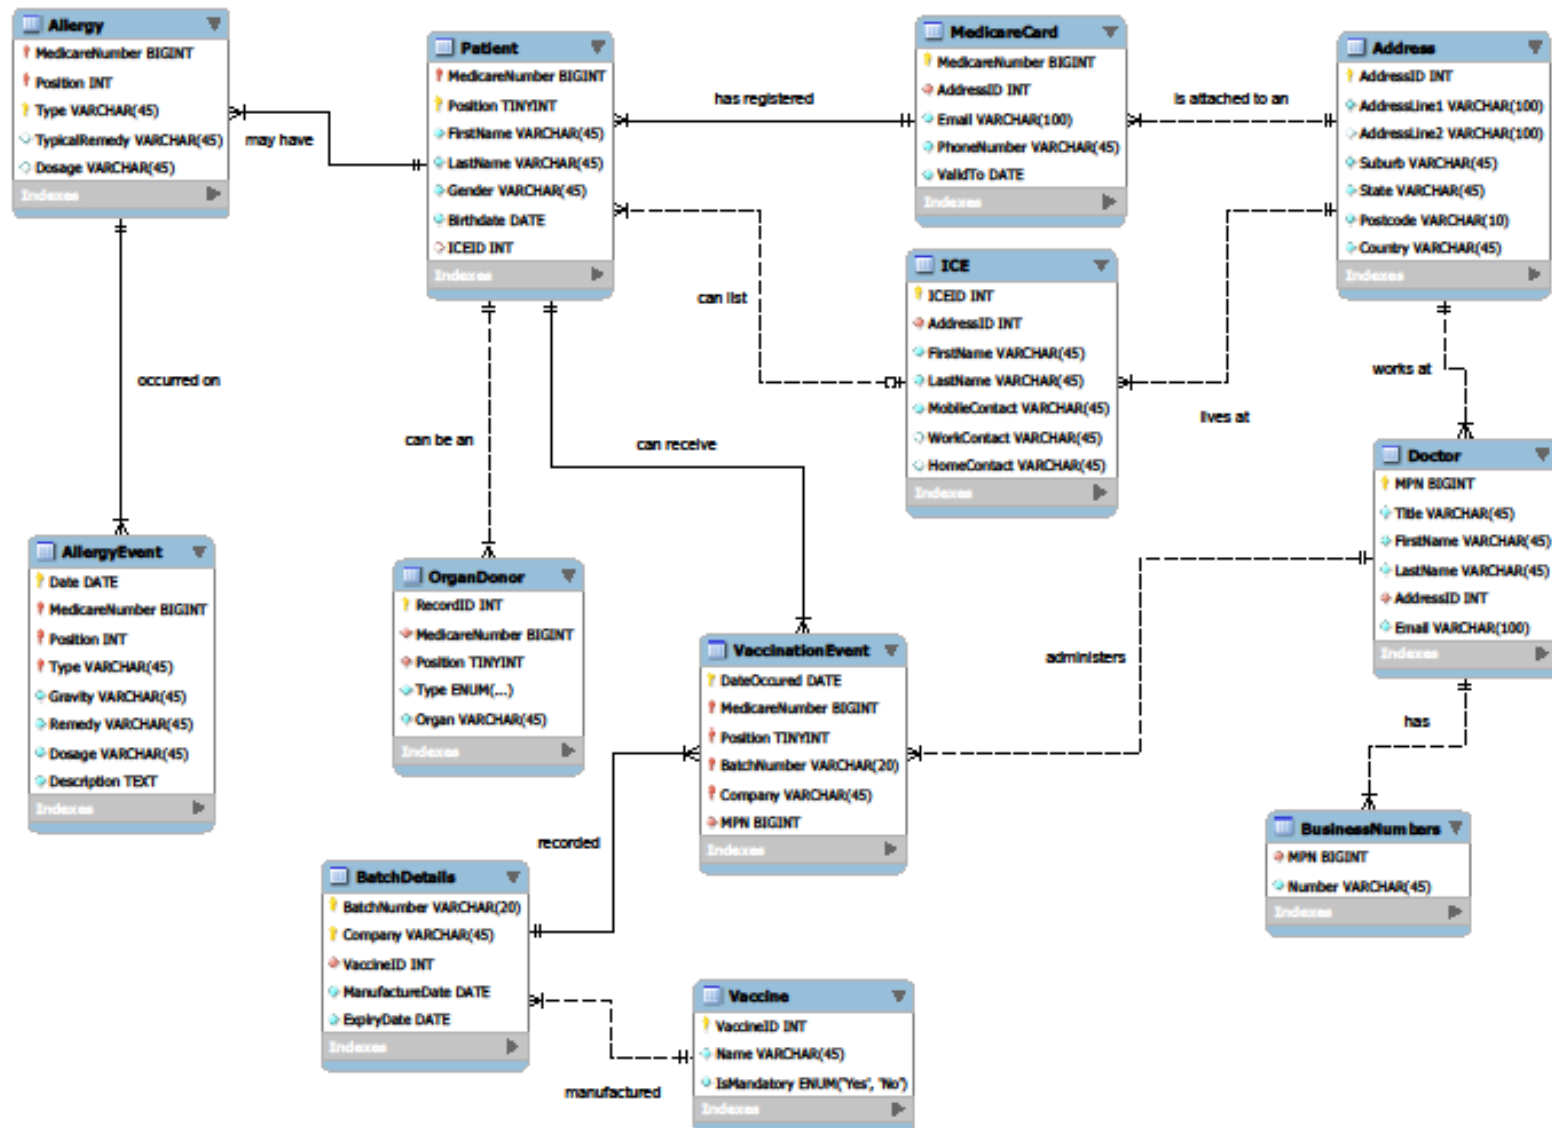
Lecture 13
Query Optimization Part I

Semester 1 2018, Week 7

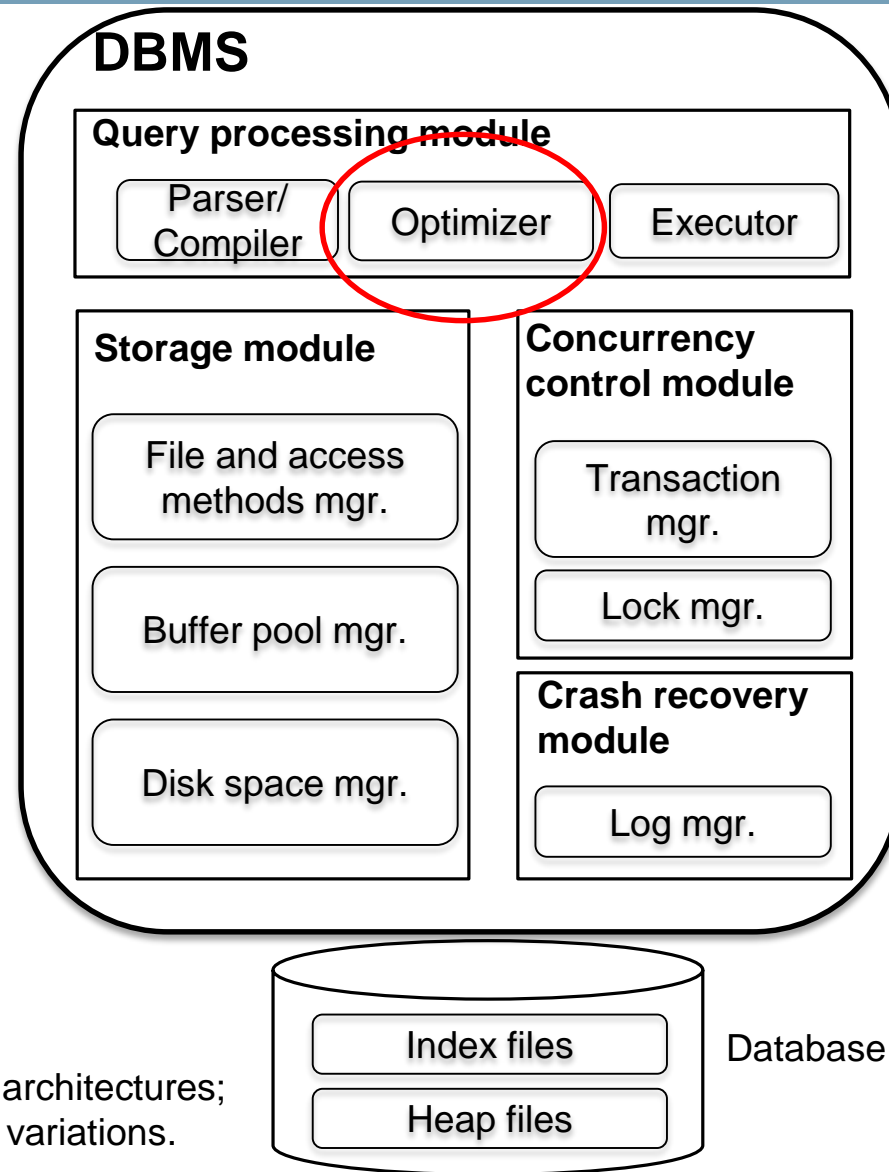


- A1 feedback is sent
 - Discuss your feedback with your tutors
 - Attend consultation hours of other tutors
 - Come to me in case of issues
- MST:
 - Wednesday, 2nd of May, 2018
 - Venue: Wilson Hall
 - Topics: (E)ER, Relational algebra, SQL

An example of a nice and elegant solution



Remember this? Components of a DBMS



**TODAY &
Next time**

This is one of several possible architectures; each system has its own slight variations.



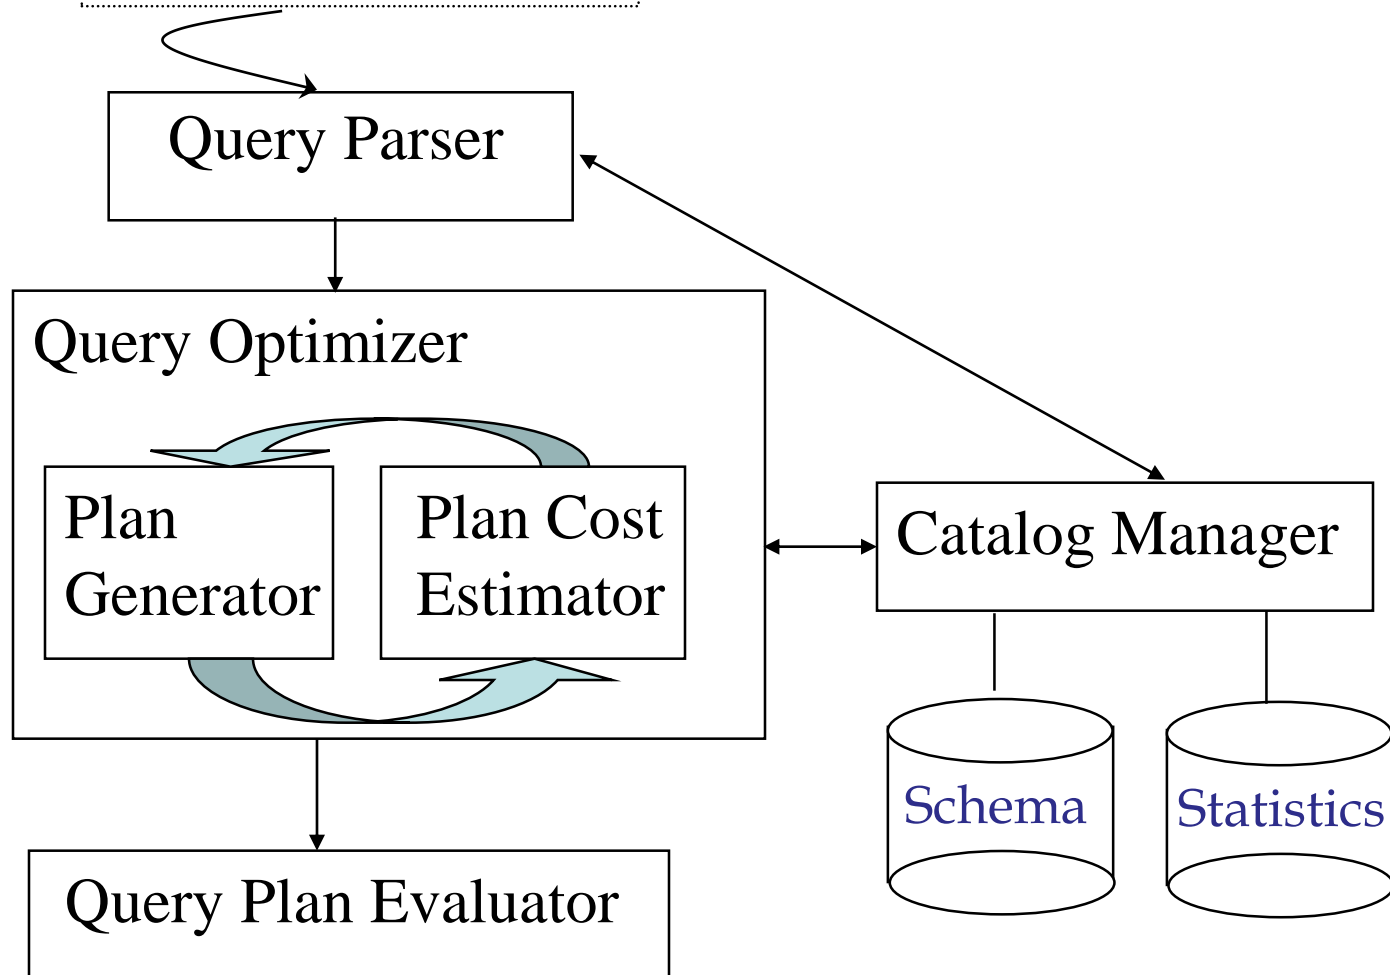
REEDUCATION

- Overview
- Query optimization
- Cost estimation

Readings: Chapter 12 and 15, Ramakrishnan & Gehrke, Database Systems

Query

```
Select *  
From Blah B  
Where B.blah = "foo"
```





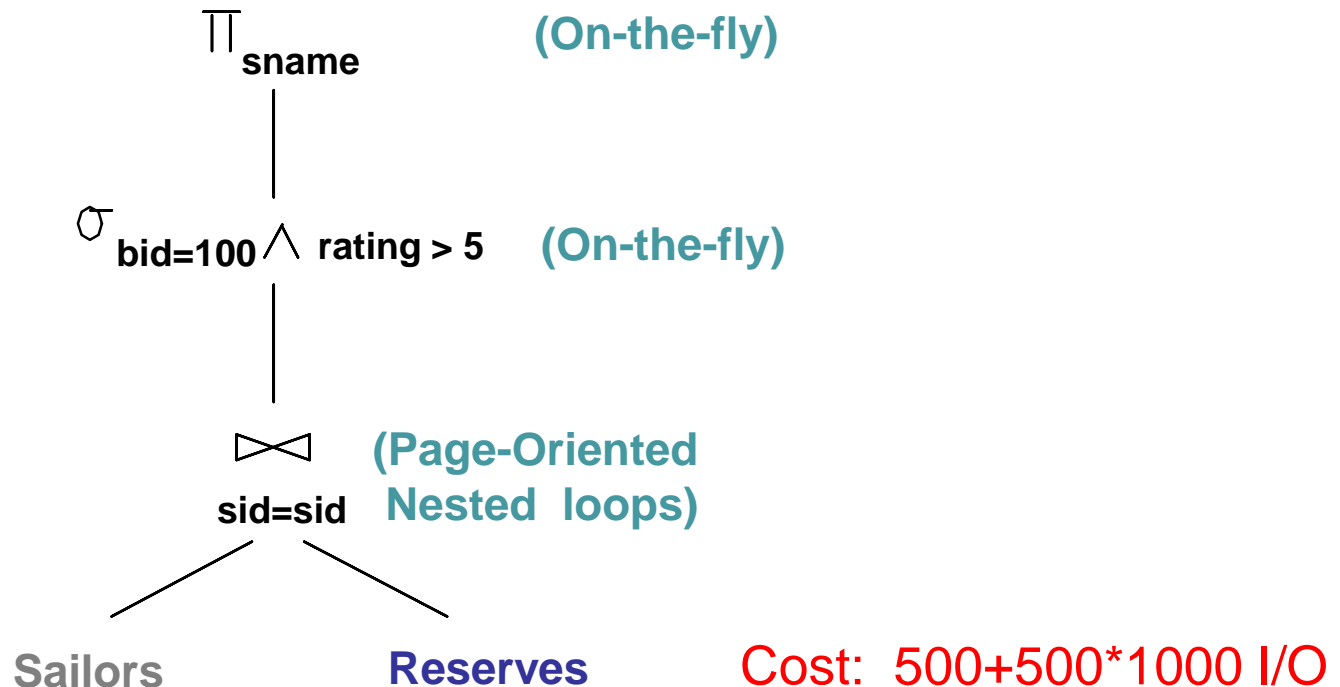
- Typically there are many ways of executing a given query, all giving the same answer
- Cost of alternative methods often **varies enormously**
- Query optimization aims to find the execution strategy with the lowest cost
- We will cover:
 - Relational algebra equivalences
 - Cost estimation
 - Result size estimation and reduction factors
 - Enumeration of alternative plans



- A tree, with relational algebra operators as nodes
- Each operator labeled with a choice of algorithm

SELECT sname from Sailors NATURAL JOIN Reserves
WHERE bid = 100 and rating > 5

Plan:





- Overview
- Query optimization
- Cost estimation

Readings: Chapter 15, Ramakrishnan & Gehrke, Database Systems



Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

Boats (*bid*: integer, *bname*: string, *color*: string)

Example:


```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

Query optimization steps:

1. Query first broken into “blocks”
2. Each block converted to relational algebra
3. Then, for each block, several alternative **query plans** are considered
4. Plan with the lowest **estimated cost** is selected

Step 1: Break query into query blocks

- Query block is any statement starting with select
- Query block = unit of optimization
- Typically inner most block is optimized first, then moving towards outers



Query:

```
SELECT S.sid  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
```

Relational algebra:

$$\pi_{S.sid}(\sigma_{B.color = \text{"red"}}(\text{Sailors} \bowtie \text{Reserves} \bowtie \text{Boats}))$$

- *Selections*: $\sigma_{c_1 \wedge \dots \wedge c_n}(R) \equiv \sigma_{c_1} \left(\dots \left(\sigma_{c_n}(R) \right) \right)$ (*Cascade*)
 $\sigma_{c_1} \left(\sigma_{c_2}(R) \right) \equiv \sigma_{c_2} \left(\sigma_{c_1}(R) \right)$ (*Commute*)
- *Projections*: $\pi_{a_1}(R) \equiv \pi_{a_1} \left(\dots \left(\pi_{a_n}(R) \right) \right)$ (*Cascade*)
 a_i is a set of attributes of R and $a_i \subseteq a_{i+1}$ for $i = 1 \dots n - 1$
- These equivalences allow us to ‘push’ selections and projections ahead of joins.

Selection:

$$\sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\text{Sailors})$$

$$\longleftrightarrow \sigma_{\text{age} < 18} (\sigma_{\text{rating} > 5} (\text{Sailors}))$$

$$\longleftrightarrow \sigma_{\text{rating} > 5} (\sigma_{\text{age} < 18} (\text{Sailors}))$$

Projection:

~~$$\pi_{\text{age}, \text{rating}} (\text{Sailors}) \longleftrightarrow \pi_{\text{age}} (\pi_{\text{rating}} (\text{Sailors}))$$~~ ?

$$\pi_{\text{age}, \text{rating}} (\text{Sailors}) \longleftrightarrow \pi_{\text{age}, \text{rating}} (\pi_{\text{age}, \text{rating}, \text{sid}} (\text{Sailors}))$$

- A projection commutes with a selection that only uses attributes retained by the projection

$$\pi_{\text{age, rating, sid}} (\sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\text{Sailors}))$$

$$\longleftrightarrow \sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\pi_{\text{age, rating, sid}} (\text{Sailors}))$$

$$\pi_{\text{age, sid}} (\sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\text{Sailors}))$$

~~$$\longleftrightarrow \sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\pi_{\text{age, sid}} (\text{Sailors}))$$~~

?

$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T \quad (\textit{Associative})$$

$$(R \bowtie S) \equiv (S \bowtie R) \quad (\textit{Commutative})$$

- These equivalences allow us to choose **different join orders**

- Converting selection + cross-product to join

$$\sigma_{S.sid = R.sid} (\text{Sailors} \times \text{Reserves})$$

$$\leftrightarrow \text{Sailors} \bowtie_{S.sid = R.sid} \text{Reserves}$$

- Selection on just attributes of S commutes with $R \bowtie S$

$$\sigma_{S.age < 18} (\text{Sailors} \bowtie_{S.sid = R.sid} \text{Reserves})$$

$$\leftrightarrow (\sigma_{S.age < 18} (\text{Sailors})) \bowtie_{S.sid = R.sid} \text{Reserves}$$

- We can also “push down” projection (*but be careful...*)

$$\pi_{S.sname} (\text{Sailors} \bowtie_{S.sid = R.sid} \text{Reserves})$$

$$\leftrightarrow \pi_{S.sname} (\pi_{sname, sid}(\text{Sailors})) \bowtie_{S.sid = R.sid} \pi_{sid}(\text{Reserves}))$$



Agenda

- Overview
- Query optimization
- Cost estimation


Readings: Chapter 15, Ramakrishnan & Gehrke, Database Systems



1. Query first broken into “blocks”
2. Each block converted to relational algebra
3. Then, for each block, several alternative **query plans** are considered
4. Plan with lowest **estimated cost** is selected

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

$\pi_{(sname)} \sigma_{(bid=100 \wedge rating > 5)} (Reserves \bowtie Sailors)$



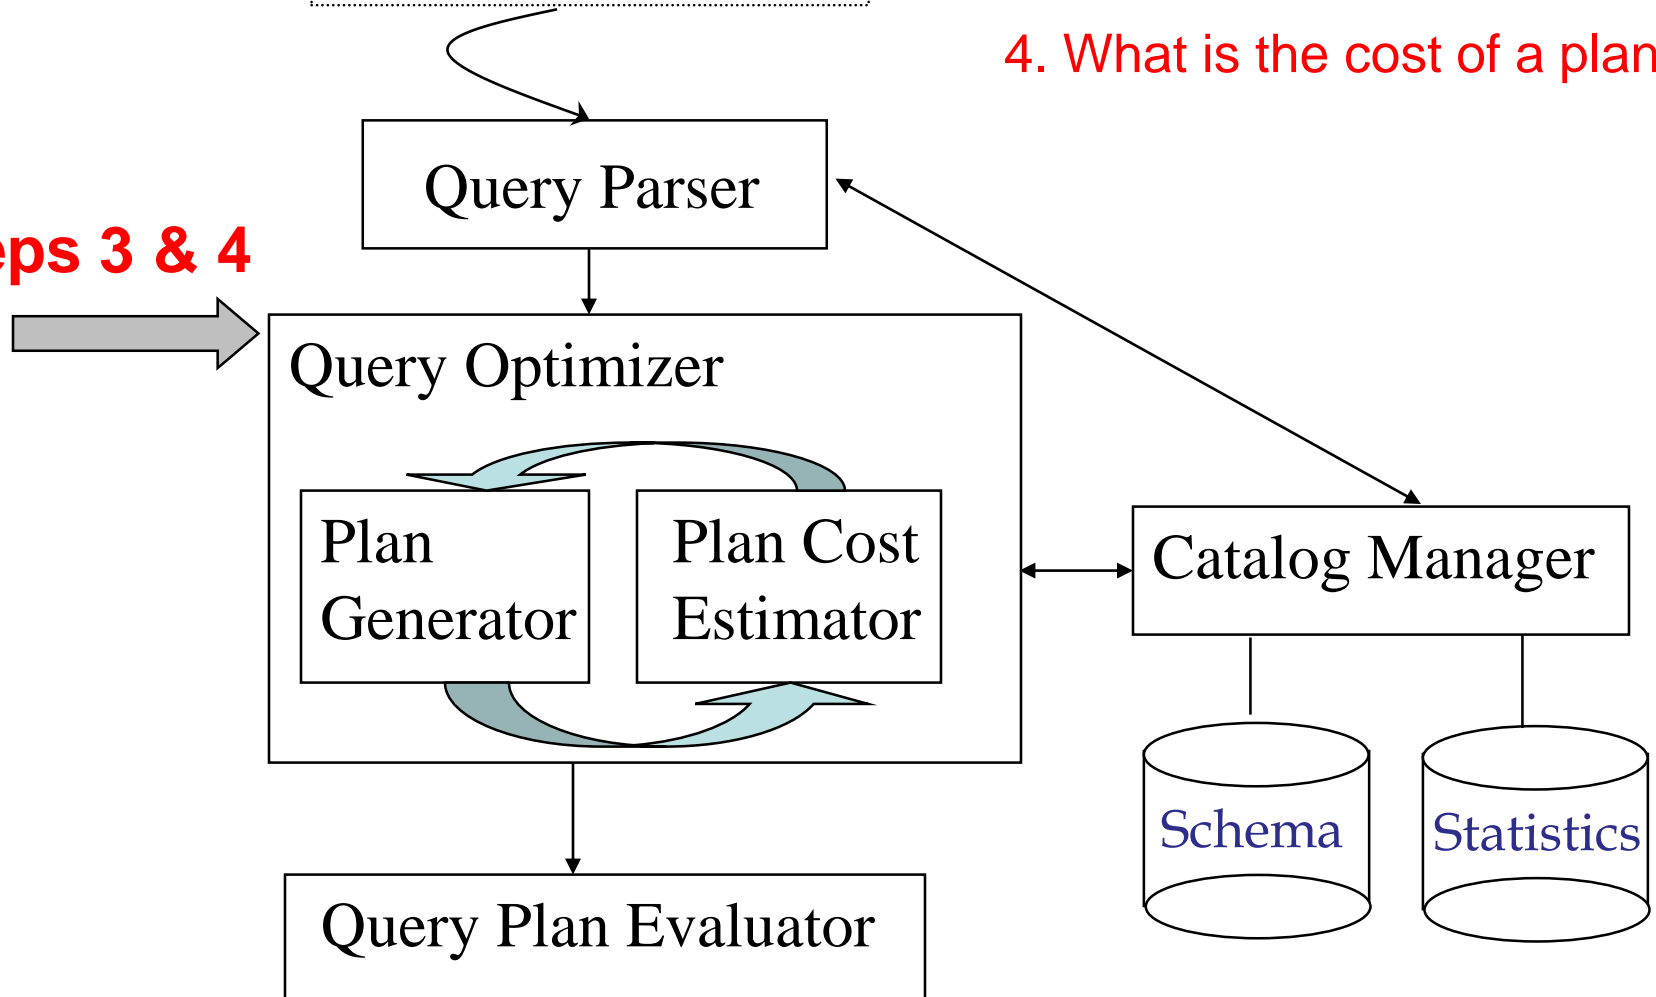


Queries

```
Select *  
From Blah B  
Where B.blah = "foo"
```

Steps 3 & 4

- 3. What plans are considered?
- 4. What is the cost of a plan?





- For each plan considered, must estimate cost:
 - Must *estimate size of result* for each operation in tree
 - Use information about input relations (from the system catalogs), and apply rules (*discussed next*)
 - Must *estimate cost* of each operation in plan tree
 - Depends on input cardinalities
 - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins)
 - *Next time we will calculate the cost of entire plans...*

- To decide on the cost, the optimizer needs information about the relations and indexes involved. This information is stored in the system **catalogs**.
- **Catalogs** typically contain at least:
 - # tuples (**NTuples**) and # pages (**NPages**) per relation
 - # distinct key values (**NKeys**) for each index (or relation attribute)
 - low/high key values (**Low/High**) for each index (or relation attribute)
 - Index height (**Height(I)**) for each tree index
 - # index pages (**NPages(I)**) for each index
- Statistics in catalogs are updated periodically

- Consider a query block:

```
SELECT attribute list
FROM relation list
WHERE predicate1 AND ... AND predicate_k
```
- Maximum number of tuples in the result is the **product** of the cardinalities of relations in the FROM clause
- Reduction factor (RF)** associated with each predicate reflects the impact of the predicate in reducing the result size. RF is also called **selectivity**.

- Single table selection:

$$\mathbf{ResultSize} = NTuples(R) \prod_{i=1..n} RF_i$$

- Joins (over k tables):

$$\mathbf{ResultSize} = \prod_{j=1..k} NTuples(R_j) \prod_{i=1..n} RF_i$$

- If there are no selections (no predicates), reduction factors are simply ignored, i.e. they are ==1



- Depend on the type of the predicate:
 1. Col = value
 $RF = 1/NKeys(Col)$
 2. Col > value
 $RF = (High(Col) - value) / (High(Col) - Low(Col))$
 3. Col < value
 $RF = (val - Low(Col)) / (High(Col) - Low(Col))$
 4. Col_A = Col_B (for joins)
 $RF = 1/ (Max (NKeys(Col_A), NKeys(Col_B)))$
 5. In no information about Nkeys or interval, use a “magic number” 1/10
 $RF = 1/10$



1. Sailors (S): NTuples(S) = 1000, Nkeys(rating) = 10 interval [1-10], age interval [0-100], Nkeys(sid)=1000

*SELECT * FROM Sailors WHERE rating = 3 AND age > 50*

NTuples(S) = 1000; RF(rating) = 1/10; RF(age) = (100-50)/(100-0)

ResultSize = NTuples(S)*1/10*((100-50)/(100-0))= 1000*0.1*0.5= 50 tuples

2. Reserves (R): NTuples(R) = 100, Nkeys(sid) = 100, Sailors per above

*SELECT * FROM Sailors as S INNER JOIN Reserves as R
ON S.SID = R.SID WHERE rating = 8 and 20 < age < 30;*

NTuples(S)= 1000; NTuples(R) = 100; RF(S.SID=R.SID)?; RF(rating)=?;

RF(20<age<30)=?

ResultSize=?



- What is query optimization/describe steps?
- Equivalence classes
- Result size estimation

- Important for Assignment 3 as well



- Query optimization Part II
 - Plan enumeration