



INFO20003 Database Systems

Dr Renata Borovica-Gajic

Lecture 14
Query Optimization Part II

Semester 1 2018, Week 7



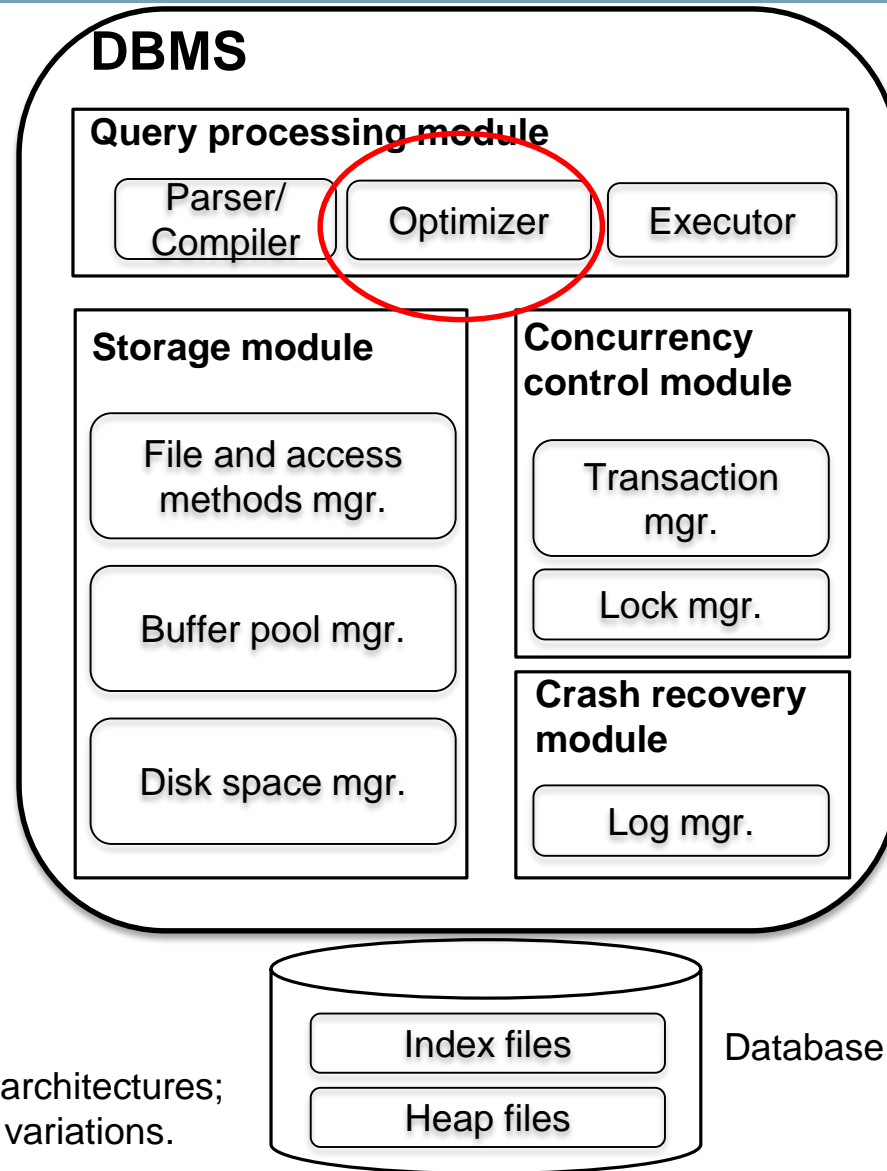
WHY ARE DATABASES COOL ?

Featuring people from industry
Part 1

SNOWFLAKE COMPUTING, USA



Remember this? Components of a DBMS



TODAY
Plan enumeration

This is one of several possible architectures; each system has its own slight variations.

Database

- When enumerating alternative plans, there are two main cases:
 - Single-relation plans
 - Multiple-relation plans (joins)
- For queries over a **single relation**:
 - Each *available* access path (file scan / index) is considered, and the one with the **lowest estimated cost** is chosen
 - Heap scan is always one alternative
 - Each index can be another alternative (if matching selection predicates)
 - Other operations can be performed on top of access paths, but they typically don't incur additional cost since they are done *on the fly* (e.g. projections, additional non-matching predicates)

1. Sequential (heap) scan of data file:

$$\text{Cost} = \text{NPages}(\mathbf{R})$$

2. Index selection over a primary key (just a single tuple):

$$\text{Cost}(\mathbf{B+Tree}) = \text{Height}(\mathbf{I}) + 1, \text{ Height is the index height}$$

$$\text{Cost}(\text{HashIndex}) = \text{ProbeCost}(\mathbf{I}) + 1, \text{ ProbeCost}(\mathbf{I}) \sim 1.2$$

3. Clustered index matching one or more predicates:

$$\text{Cost}(\mathbf{B+Tree}) = (\text{NPages}(\mathbf{I}) + \text{NPages}(\mathbf{R})) * \prod_{i=1..n} RF_i$$

$$\text{Cost}(\text{HashIndex}) = \text{NTuples}(\mathbf{R}) * \prod_{i=1..n} RF_i * 2.2$$

4. Non-clustered index matching one or more predicates:

$$\text{Cost}(\mathbf{B+Tree}) = (\text{NPages}(\mathbf{I}) + \text{NTuples}(\mathbf{R})) * \prod_{i=1..n} RF_i$$

$$\text{Cost}(\text{HashIndex}) = \text{NTuples}(\mathbf{R}) * \prod_{i=1..n} RF_i * 2.2$$

Let's say that Sailors(S) has 500 pages, 40000 tuples, $NKeys(rating) = 10$

```
SELECT S.sid FROM Sailors S WHERE S.rating=8
```

• Result size = $(1/NKeys(I)) * NTuples(S) = (1/10) * 40000 = 4000$ tuples

1. If we have $I(rating)$, $NPages(I) = 50$:

– Clustered index:

Cost = $(1/NKeys(I)) * (NPages(I) + NPages(S)) = (1/10) * (50 + 500) = 55$ I/O

– Unclustered index:

Cost = $(1/NKeys(I)) * (NPages(I) + NTuples(S)) = (1/10) * (50 + 40000) = 4005$ I/O

2. If we have an $I(sid)$, $NPages(I) = 50$:

– Cost = ?, Result size = ?

– Would have to retrieve all tuples/pages. With a clustered index, the cost is $50 + 500$, with unclustered index, $50 + 40000$

3. Doing a file scan:

– Cost = $NPages(S) = 500$

Steps:


1. Select **order** of relations
 - E.g. $SxRxB$, or $SxBxR$ or $RxSxB...$
 - maximum possible orderings = $N!$
2. For each join, select **join algorithm**
 - E.g. Hash join, Sort-Merge Join...
3. For each input relation, select access method
 - Heap Scan, or various index alternatives

Q: How many plans are there for a query over N relations?

Back-of-envelope calculation:

- With 3 join algorithms, I indexes per relation:
plans $\approx [N!] * [3^{(N-1)}] * [(I + 1)^N]$
- Suppose $N = 3$, $I = 2$: # plans $\approx 3! * 3^2 * 3^3 = 1458$ plans
- This is just for illustration – you don't need to remember this

- As number of joins increases, number of alternative plans grows rapidly → *need to restrict search space*
- Fundamental decision in System R (first DBMS): **only left-deep join trees** are considered
 - Left-deep trees allow us to generate all *fully pipelined* plans
 - Intermediate results are not written to temporary files









```
SELECT S.sname, B.bname, R.day  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid
```

- Let's assume:
 - Two join algorithms to choose from:
 - Hash-Join
 - NL-Join (page-oriented)
 - Unclustered B+Tree index: $I(R.sid)$; $NPages(I) = 50$
 - No other indexes
 - S: $NPages(S) = 500$, $NTuplesPerPage(S) = 80$
 - R: $NPages(R) = 1000$, $NTuplesPerPage(R) = 100$
 - B: $NPages(B) = 10$
 - 100 R \bowtie S tuples fit on a page

Candidate Plans

```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```


1. Enumerate relation orderings:




*** Prune plans with cross-products immediately!**

```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```

2. Enumerate join algorithm choices:




+ do the same
for other plans





```
SELECT S.sname, B.bname, R.day
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
```


3. Enumerate **access method** choices:



+ do same for
other plans



Now estimate the cost of each plan



S: NPages(S) = 500, NTuplesPerPage(S) = 80
 R: NPages(R) = 1000, NTuplesPerPage(R) = 100
 B: NPages(B) = 10
 100 R \bowtie S tuples fit on a page
 All 3 relations are Heap Scan

Calculating cost:

SxR

Cost (SxR) = 500 + 500*1000 = 500500

(SxR)xB

Result size (SxR) = 100000*40000 * $\frac{1}{40000}$ = 100000 tuples = 1000 pages


Cost(xB) = ~~1000~~ + 1000*10 = 10000

Already read – left deep plans apply pipelining

Total Cost = 500 + 500*1000 + 1000 * 10 = 510500 I/O



Now estimate the cost of each plan



S: NPages(S) = 500, NTuplesPerPage(S)= 80
 R: NPages(R) = 1000, NTuplesPerPage(R) = 100
 B: NPages(B) = 10
 100 R \bowtie S tuples fit on a page
 All 3 relations are Heap Scan

Calculating cost:

SxR

$$\text{Cost (SxR)} = 500 + 500 \times 1000 = 500500$$

(SxR)xB


Result size (SxR) = $100000 \times 40000 \times \frac{1}{40000} = 100000$ tuples = 1000 pages

$$\text{Cost(xB)} = 3 \times 1000 + 3 \times 10 = 2 \times 1000 + 3 \times 10 = 2030$$

Already read once – left deep plans apply pipelining


$$\text{Total Cost} = 500 + 500 \times 1000 + 2 \times 1000 + 3 \times 10 = 502530 \text{ I/O}$$

Plan 3:



S: $NPages(S) = 500$, $NTuplesPerPage(S) = 80$
R: $NPages(R) = 1000$, $NTuplesPerPage(R) = 100$
B: $NPages(B) = 10$
 $100 R \bowtie S$ tuples fit on a page
All 3 relations are Heap Scan

Plan 4:






Calculating cost:

Cost (P3) = ?

Cost (P4) = ?

SMJ : 2 passes, RxB: 10 tuples per page





- Understand plan enumeration and cost various plans
- Important for Assignment 3 as well



- Normalization