



COMP20008 Elements of Data Processing

Semester 2 2018

Lecture 5: Recommender Systems



THE UNIVERSITY OF
MELBOURNE

Plan today

- Complete section on outlier detection
- Recommender systems and collaborative filtering
- Types of similarity for imputation of missing values
 - Item-Item
 - User-User
 - Matrix factorisation
- Question to consider during lecture: Are we doing cleaning or prediction?



THE UNIVERSITY OF
MELBOURNE

Recommender systems: missing data

- Movie Recommender systems

Person	Star Wars	Batman	Jurassic World	The Martian	The Revenant	Lego Movie	Selma	...
James	3	2	-	-	-	1	-	
John	-	-	1	2	-	-	-	
Jill	1	-	-	3	2	1	-	

Users and movies

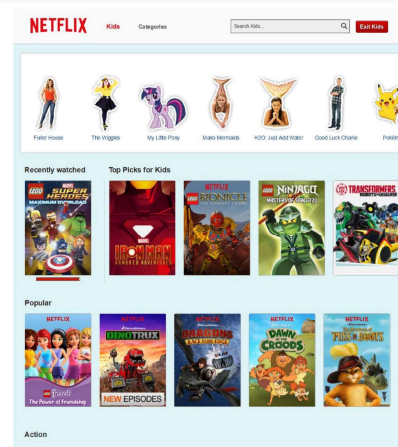
Each user only rates a few movies (say 1%)

Netflix wants to predict the missing ratings for each user



THE UNIVERSITY OF
MELBOURNE

Netflix






The Revenant: A Novel of
Revenge
› Michael Punke
1,250
Paperback
\$9.52



Ready Player One: A Novel
› Ernest Cline
9,210
Paperback
\$8.37



The Life We Bury
› Allen Eskens
1,896
Paperback
\$8.76




The 5th Wave: The First
Book of the 5th Wave
Series
› Rick Yancey
2,006
Paperback
\$6.70

- “75% of what people watch is from some sort of recommendation” (Netflix)
- “If I have 3 million customers on the web, I should have 3 million stores on the web.” (Amazon CEO)

- IMDb
- Online dating
- Twitter: “Who to Follow”, what to retweet
- Spotify, youtube: music recommendation
- LinkedIn/Facebook: who to add as a contact, jobs of interest, news of interest
- Tourist attraction apps
- University subjects ... ? Subject discussion forums ... ?

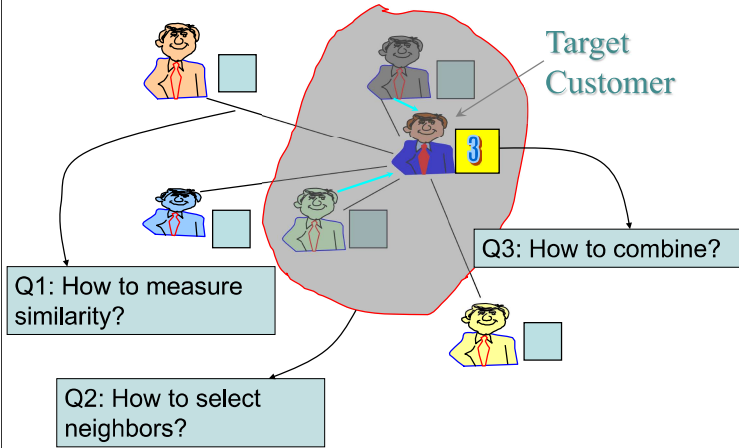
- Each user has a profile
- Users rate items
 - Explicitly: Give a score
 - Implicitly: web usage mining
 - Time spent in viewing the item
 - Navigation path
 - Etc...
- System does the rest, How?

- *Collaborative Filtering: Make predictions about a user's missing data according to the behaviour of many other users*
 - Look at users **collective** behavior
 - Look at the active user **history**
 - Combine!



- User based methods
 - Identify like-minded users
- Item based methods
 - Identify similar items
- Model (matrix) based methods
 - Solve an optimization problem and identify latent factors

	Items					
	Item1	Item2	Item3	Item4	Item5	Item6
User1	17	-	20	18	17	18.5
User2	8	-	????	17	14	17.5
User3	-	-	17	18	18.5	17.5
User4	-	-	-	18	17.5	18
User5	17	-	18	19	15.5	-
User6	-	-	17.5	-	16	-
User7	15	17.5	-	17	-	17
User8	18	-	-	-	17	16.5
User9	18	17	-	-	18.5	17
User10	19	17	-	-	-	16.5
User11	17	18.5	19	19	-	-
User12	14	19	17	-	-	15.5
User13	-	16	-	-	17	-
User14	20	18.5	-	18	-	18



U1	17	-	20	18	17	18.5
U2	8	-	-	17	14	17.5

	i_1	i_n
u_1		
u_2		

$$SIM(U1, U2) = ((17 - 8)^2 + (18.1 - 14.1)^2 + (20 - 14.1)^2 + (18 - 17)^2 + (17 - 14)^2 + (18.5 - 17.5)^2)$$

- Compute mean value for User1's missing values (18.1)
- Compute mean value for User2's missing values (14.1)
- Compute squared Euclidean distance between resulting vectors

							i_1	i_n
User1	17	-	20	18	17	18.5	u_1	
User2	8	-	-	17	14	17.5	u_2	

$$Sim(User1, User2) = \frac{6}{6-2} ((17-8)^2 + (18-17)^2 + (17-14)^2 + (18.5-17.5)^2)$$

- Compute squared Euclidean distance between vectors, summing only pairs without missing values
- 2 out of the 6 pairs have at least one missing value
- Scale the result, according to percentage of pairs with a missing value


User1	12	2.5	20	-	17	-	3.5
User2	13	-	-	17	14	17.5	4.5

Using Method 2, $SIM(User1, User2) = ?$



$$SIM(User1, User2) = \frac{7}{3} (|12 - 13|^2 + |17 - 14|^2 + |3.5 - 4.5|^2)$$

$$= \frac{7}{3} (1 + 9 + 1) = 25.66$$

- Instead of Euclidean distance can also use other measures to assess similarity, e.g.
 - Correlation (we will look at later in subject)
 - Cosine similarity (angle between user profile vectors)




- At runtime
 - Need to *select* users to compare to
 - Could choose the top-k most similar users
 - *Combining*: Prediction of rating is the (weighted) average of the values from the top-k similar users
- Can make more efficient by computing clusters of users offline
 - At runtime find nearest cluster and use the centre of the cluster as the rating prediction
 - Faster (more scalable) but a little less accurate



- Achieve good quality in practice
- The more processing we push offline, the better the method scale
- However:
 - User preference is **dynamic**
 - High update frequency of offline-calculated information
 - No recommendation for **new users**
 - We don't know much about them yet

- Search for similarities among items
- All computations can be **done offline**
- Item-Item similarity is more stable than user-user similarity
 - No need for frequent updates:

- Same as in user-user similarity but on item vectors
 - Find similar items to the one whose rating is missing
 - E.g. For item i_i compute its similarity to each other item i_j



- Offline phase. For each item
 - Determine its k-most similar items
 - Can use same type of similarity as for user-based
- Online phase:
 - Predict rating r_{aj} for a given user-item pair as a weighted sum over k-most similar items that they rated

$$r_{aj} = \frac{\sum_{i \in \text{k-similar items}} \text{sim}(i, j) \times r_{ai}}{\sum_{i \in \text{k-similar items}} \text{sim}(i, j)}$$

User a	8		r_{aj}	9	15
	Item j				

Practice example

Users	Titanic	Batman	Inception	Superman	The Martian	Jurassic World
Michelle	2.5		3	3.5	2.5	3
Tom	3	3.5		5	3	3.5
Lao	2.5	3		3.5		4
Chan		3.5	3	4	2.5	
Mary		4	2	3	2	3
Tim	3	4	?	5	3.5	3
John		4.5		4	1	

Item j : Inception
User a : Tim

Offline phase: we calculate k -most similar items for item j . Let's say $k = 3$.

Similarity	Titanic	Batman	Superman	The Martian	Jurassic World
Inception-Method 1	1.06	3.28	3.95	2.04	2.02
Inception-Method 2	3.5	7.22	3.5	1.65	3.5

Practice example cont.

Users	Titanic	Batman	Inception	Superman	The Martian	Jurassic World
Michelle	2.5		3	3.5	2.5	3
Tom	3	3.5		5	3	3.5
Lao	2.5	3		3.5		4
Chan		3.5	3	4	2.5	
Mary		4	2	3	2	3
Tim	3	4	?	5	3.5	3
John		4.5		4	1	

Item j : Inception
User a : Tim

$$r_{aj} = \frac{\sum_{i \in k\text{-similar items}} \text{sim}(i, j) \times r_{ai}}{\sum_{i \in k\text{-similar items}} \text{sim}(i, j)}$$

Similarity	Superman	The Martian	Jurassic World
$\text{Sim}(i, j)$	3.5	1.65	3.5

Online phase:

$$r_{aj} = \frac{3.5 \times 5 + 1.65 \times 3.5 + 3.5 \times 3}{3.5 + 1.65 + 3.5} = 3.9$$

Matrix based techniques

- Treat the User-Item Rating table R as a matrix
 - Use matrix factorisation of this Rating Table

Matrix based techniques: Rating table R

	Items					
	Item1	Item2	Item3	Item4	Item5	Item6
User1	17	-	20	18	17	18.5
User2	8	-	-	17	14	17.5
User3	-	-	17	18	18.5	17.5
User4	-	-	-	18	17.5	18
User5	17	-	18	19	15.5	-
User6	-	-	17.5	-	16	-
User7	15	17.5	-	17	-	17
User8	18	-	-	-	17	16.5
User9	18	17	-	-	18.5	17
User10	19	17	-	-	-	16.5
User11	17	18.5	19	19	-	-
User12	14	19	17	-	-	15.5
User13	-	16	-	-	17	-
User14	20	18.5	-	18	-	18

- We are familiar with factorisation of numbers
 $15=3*5$
 $99=3*33$
 $1000=10*100$

We can also do approximate factorisation
 $17 \approx 6*2.8$ (RHS= 16.8, an error of 0.2)
 $167 \approx 17*9.8$ (RHS=166.6, an error of 0.4)

Given a matrix R, we can find matrices U and V such that when U and V are multiplied together

$$R \approx UV$$

- R is $m \times n$, U is $m \times k$ and V is $k \times n$
 - k is the "number of latent factors"

For example, suppose $R = \begin{bmatrix} 5 & 2 & 3 & 6 \\ 4 & 4 & 6 & 11 \\ 3 & 19 & 2 & 7 \\ 3 & 8.5 & 4 & 2 \end{bmatrix}$
R is a 4×4 matrix

$$\begin{bmatrix} 5 & 2 & 3 & 6 \\ 4 & 4 & 6 & 11 \\ 3 & 19 & 2 & 7 \\ 3 & 8.5 & 4 & 2 \end{bmatrix} \approx \begin{bmatrix} 0.34776 & 1.97802 \\ 0.71609 & 3.13615 \\ 4.27876 & 0.58287 \\ 1.88074 & 0.56923 \end{bmatrix} \begin{bmatrix} 0.58367 & 4.40189 & 0.44605 & 1.04492 \\ 1.52915 & 0.26346 & 1.75046 & 3.09976 \end{bmatrix}$$

$$= \begin{bmatrix} 3.22769 & 2.05196 & 3.61758 & 6.49480 \\ 5.21363 & 3.97844 & 5.80912 & 10.46959 \\ 3.3887 & 18.98823 & 2.92886 & 6.27777 \\ 1.96819 & 8.42882 & 1.83534 & 3.72973 \end{bmatrix}$$

We can compute the error (squared distance between R and UV). The smaller it is, the better the fit of the factorisation.

$$(5 - 3.22769)^2 + (2 - 2.05196)^2 + (3 - 3.61758)^2 + \dots$$

$$(4 - 1.83534)^2 + (2 - 3.72973)^2$$

- Details of how to compute the matrix factorisation are beyond the scope of our study.
- Intuitively, factorisation algorithms search over lots of choices for U and V, with the aim of making the error as low as possible
- If there are missing values in R, ignore these when computing the error.

$$\begin{bmatrix} 5 & - & - & 6 \\ - & 4 & 6 & 11 \\ - & 19 & 2 & 7 \\ 3 & 8.5 & - & - \end{bmatrix} \approx \begin{bmatrix} 1.51261 & 1.65457 \\ -0.0474 & 3.56317 \\ 3.88351 & 1.50482 \\ 1.76637 & 0.56005 \end{bmatrix} \begin{bmatrix} 1.07179 & 4.42771 & -0.13516 & 0.60378 \\ 2.01538 & 1.18272 & 1.67926 & 3.08647 \end{bmatrix}$$

$$= \begin{bmatrix} 4.95572 & 8.65430 & 2.57402 & 6.02008 \\ 7.13025 & 4.00394 & 5.98995 & 10.96899 \\ 7.19512 & 18.97488 & 2.00210 & 6.98942 \\ 3.02190 & 8.48338 & 0.70173 & 2.79509 \end{bmatrix}$$

$$\text{Error} = (5 - 4.95572)^2 + (6 - 6.02008)^2 + (4 - 4.00394)^2 + (6 - 5.98995)^2 + \dots$$

The product of the two factors U and V, has no missing values. We can use this to predict our missed entries.
E.g. $R_{12}=8.65430$

	Items					
	Item1	Item2	Item3	Item4	Item5	Item6
User1	17	-	20	18	17	18.5
User2	8	-	13.48	17	14	17.5
User3	-	-	17	18	18.5	17.5
User4	-	-	-	18	17.5	18
User5	17	-	18	19	15.5	-
User6	-	-	17.5	-	16	-
User7	15	17.5	-	17	-	17
User8	18	-	-	-	17	16.5
User9	18	17	-	-	18.5	17
User10	19	17	-	-	-	16.5
User11	17	18.5	19	19	-	-
User12	14	19	17	-	-	15.5
User13	-	16	-	-	17	-
User14	20	18.5	-	18	-	18

- Real answer for (User 2, Item 3) is 13.5
 - Matrix technique predicts 13.48. Low error for this cell.
- Real answer for (User 13, Item 1) is 17.
 - Matrix technique predicts 15.3. Error is a little higher for this cell.
- In general, the prediction error varies across the cells, but taking all missing cells as a whole, the method aims to make predictions with low average error

- Commercial recommender systems (Netflix, Amazon) use variations of matrix factorisation.
- In 2009, Netflix offered a prize of \$USD 1,000,000 in a competition to see which algorithms were most effective for predicting user-movie ratings.
 - Anonymised training data released to public: 100 million ratings by 480k users of 17.8k movies
 - Won by "BellKor's Pragmatic Chaos" team
- A follow up competition was cancelled due to privacy concerns ... [We will elaborate when we get to topic on privacy]

- Many challenging issues in deployment of recommendations
 - Interpretability of recommendations?
 - How to be fair to rare items?
 - How to avoid only recommending popular items?
 - How to handle new users?

- See
 - Matrix Factorization Techniques for Recommender Systems. Koren, Bell and Volinsky. IEEE Xplore, Vol 42, 2009. Available on the LMS in Week 3 section.
- Some slides based on “Data Mining Concepts and Techniques”, Han et al, 2nd edition 2006.