

# INFORMED SEARCH ALGORITHMS

## CHAPTER 3, SECTIONS 5–6

# Outline

- ◇ Best-first search
- ◇  $A^*$  search
- ◇ Heuristics
- ◇ Hill-climbing

## Review: General search

```
function GENERAL-SEARCH(problem, QUEUING-FN) returns a solution, or failure
  nodes ← MAKE-QUEUE(MAKE-NODE(INITIAL-STATE[problem]))
  loop do
    if nodes is empty then return failure
    node ← REMOVE-FRONT(nodes)
    if GOAL-TEST[problem] applied to STATE(node) succeeds then return node
    nodes ← QUEUING-FN(nodes, EXPAND(node, OPERATORS[problem]))
  end
```

A strategy is defined by picking the *order of node expansion*

## Best-first search

Idea: use an *evaluation function* for each node  
– estimate of “desirability”

⇒ Expand most desirable unexpanded node

Implementation:


QUEUEINGFN = insert successors in decreasing order of desirability

Special cases:

greedy search

A\* search

# Romania with step costs in km



Straight-line distance  
to Bucharest

<b>Arad</b>	366
<b>Bucharest</b>	0
<b>Craiova</b>	160
<b>Dobreta</b>	242
<b>Eforie</b>	161
<b>Fagaras</b>	178
<b>Giurgiu</b>	77
<b>Hirsova</b>	151
<b>Iasi</b>	226
<b>Lugoj</b>	244
<b>Mehadia</b>	241
<b>Neamt</b>	234
<b>Oradea</b>	380
<b>Pitesti</b>	98
<b>Rimnicu Vilcea</b>	193
<b>Sibiu</b>	253
<b>Timisoara</b>	329
<b>Urziceni</b>	80
<b>Vaslui</b>	199
<b>Zerind</b>	374


## Greedy search

Evaluation function  $h(n)$  (heuristic)  
= estimate of cost from  $n$  to *goal*


E.g.,  $h_{\text{SLD}}(n)$  = straight-line distance from  $n$  to Bucharest

Greedy search expands the node that *appears* to be closest to goal

## Greedy search example




## Greedy search example

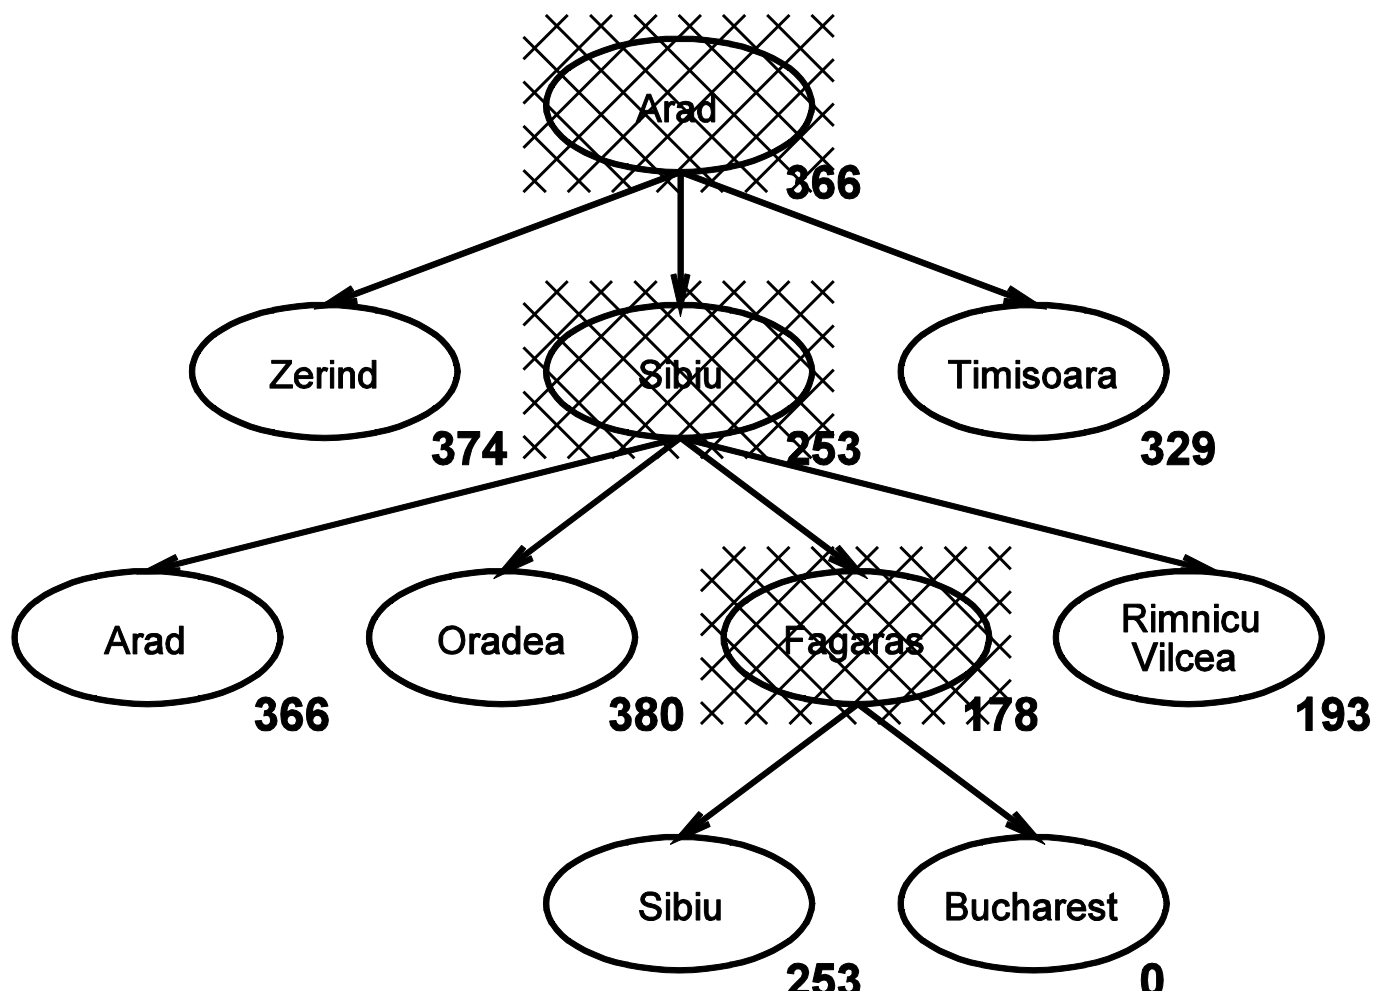




## Greedy search example



# Greedy search example



# Properties of greedy search

Complete??

Time??

Space??

Optimal??

## Properties of greedy search

Complete?? No – can get stuck in loops, e.g., Iasi to Fagaras

Iasi  $\rightarrow$  Neamt  $\rightarrow$  Iasi  $\rightarrow$  Neamt  $\rightarrow$

Complete in finite space with repeated-state checking

Time??  $O(b^m)$ , but a good heuristic can give dramatic improvement

Space??  $O(b^m)$ —keeps all nodes in memory

Optimal?? No

## A\* search

Idea: avoid expanding paths that are already expensive

Evaluation function  $f(n) = g(n) + h(n)$

$g(n)$  = cost so far to reach  $n$  (path cost)

$h(n)$  = estimated cost to goal from  $n$

$f(n)$  = estimated total cost of path through  $n$  to goal

A\* search uses an *admissible* heuristic

i.e.,  $h(n) \leq h^*(n)$  where  $h^*(n)$  is the *true* cost from  $n$ .


E.g.,  $h_{\text{SLD}}(n)$  never overestimates the actual road distance

Theorem: A\* search is optimal


# A\* search example

Arad  
**366**

# A\* search example

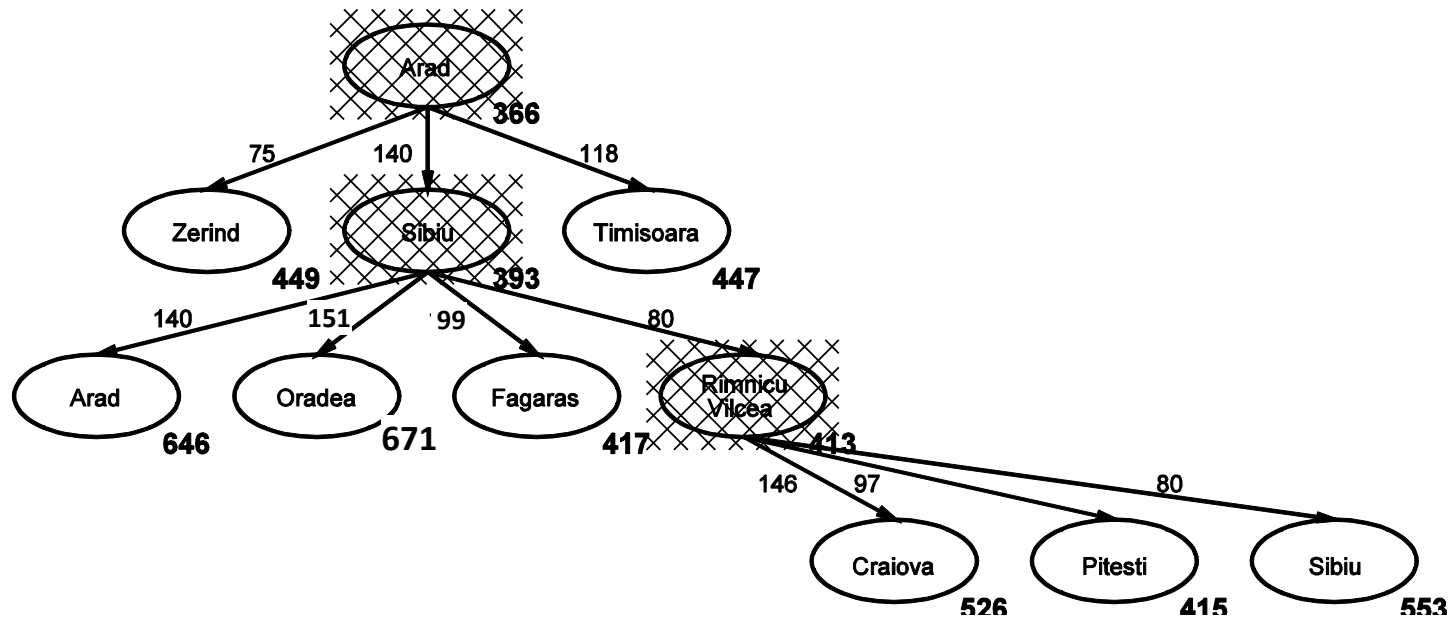


# A\* search example

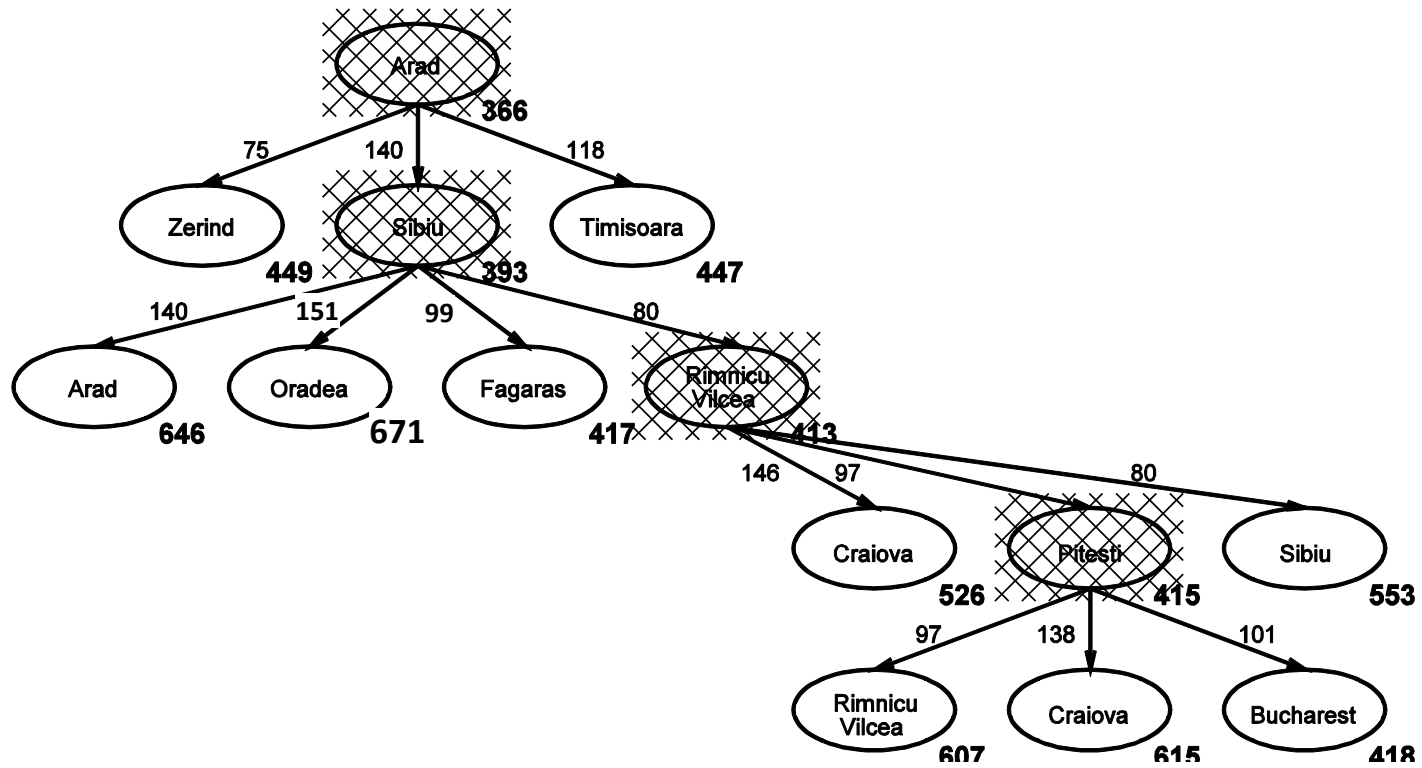




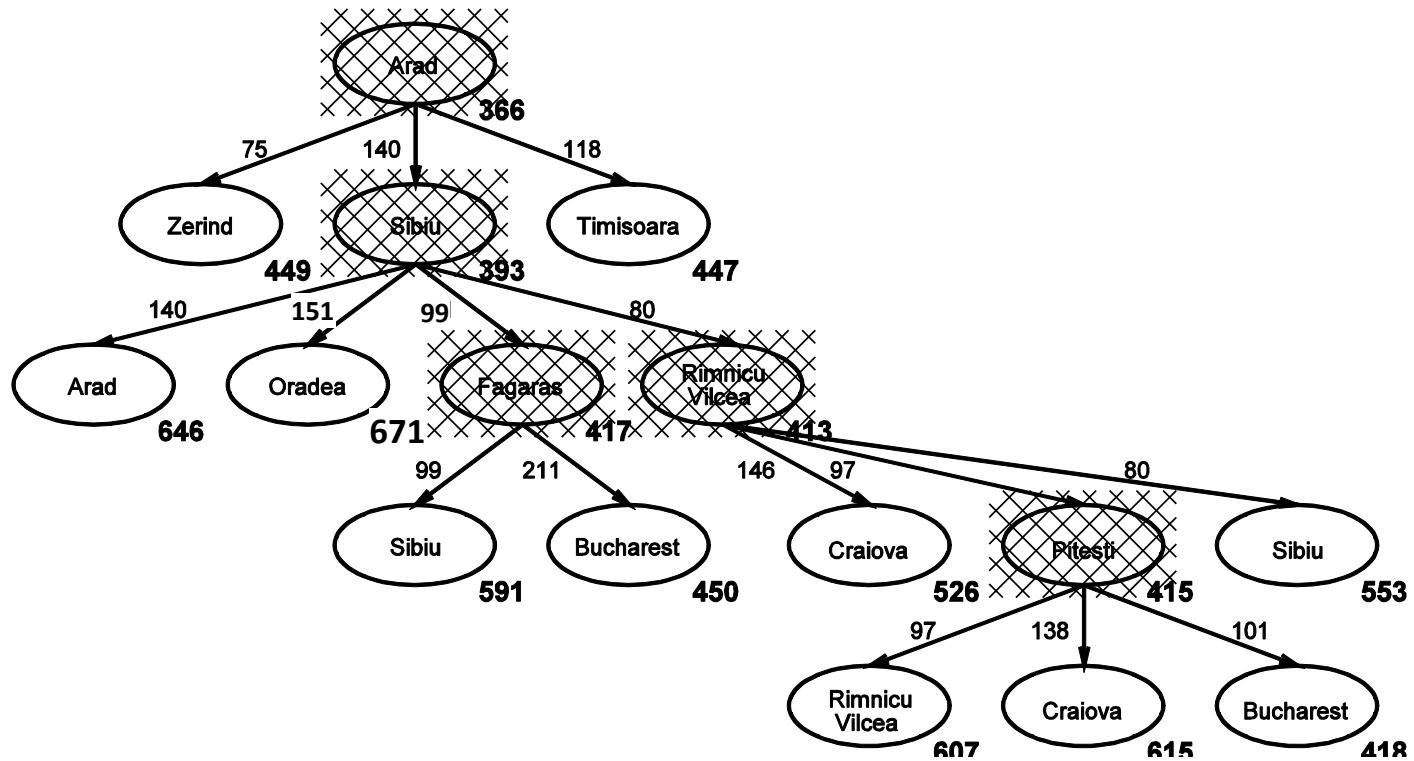
# A\* search example



# A\* search example




# A\* search example



## Optimality of $A^*$ (standard proof)

Suppose some suboptimal goal  $G_2$  has been generated and is in the queue. Let  $n$  be an unexpanded node on a shortest path to an optimal goal  $G$ .



$$\begin{aligned} f(G_2) &= g(G_2) && \text{since } h(G_2) = 0 \\ &> g(G) && \text{since } G_2 \text{ is suboptimal} \\ &\geq f(n) && \text{since } h \text{ is admissible} \end{aligned}$$


Since  $f(G_2) > f(n)$ ,  $A^*$  will never select  $G_2$  for expansion

# Optimality of $A^*$ (more useful)

Lemma:  $A^*$  expands nodes in order of increasing  $f$  value

Gradually adds “ $f$ -contours” of nodes (cf. breadth-first adds layers)

Contour  $i$  has all nodes with  $f = f_i$ , where  $f_i < f_{i+1}$



## Properties of $A^*$

Complete?? Yes, unless there are infinitely many nodes with  $f \leq f(G)$

Time?? Exponential in [relative error in  $h \times$  length of soln.]

Space?? Keeps all nodes in memory

Optimal?? Yes—cannot expand  $f_{i+1}$  until  $f_i$  is finished

## The heuristic can control $A^*$ 's behaviour

- ◇ If  $h(n)$  is very high relative to  $g(n)$ , then only  $h(n)$  plays a role, and  $A^*$  turns into ... Search.
- ◇ If  $h(n)$  is 0, then only  $g(n)$  plays a role, and  $A^*$  turns into ... Search, which finds the optimal solution.
- ◇ If  $h(n)$  is ... than the cost of moving from  $n$  to the goal, then  $A^*$  is guaranteed to find the shortest path. The ...  $h(n)$  is, the more node  $A^*$  expands, making it slower.
- ◇ If  $h(n)$  is ... than the cost of moving from  $n$  to the goal, then  $A^*$  is not guaranteed to find a shortest path, but it can run faster.
- ◇ If  $h(n)$  is ... to the cost of moving from  $n$  to the goal, then  $A^*$  will only follow the best path and never expand anything else, making it very fast.

## The heuristic can control A\*'s behaviour


- ◇ If  $h(n)$  is very high relative to  $g(n)$ , then only  $h(n)$  plays a role, and A\* turns into Greedy Best-First-Search.
- ◇ If  $h(n)$  is 0, then only  $g(n)$  plays a role, and A\* turns into Uniform Cost Search, which finds the optimal solution.
- ◇ If  $h(n)$  is always lower than (or equal to) the cost of moving from  $n$  to the goal, then A\* is guaranteed to find a shortest path. The lower  $h(n)$  is, the more node A\* expands, making it slower.
- ◇ If  $h(n)$  is sometimes greater than the cost of moving from  $n$  to the goal, then A\* is not guaranteed to find a shortest path, but it can run faster.
- ◇ If  $h(n)$  is exactly equal to the cost of moving from  $n$  to the goal, then A\* will only follow the best path and never expand anything else, making it very fast.




## Examples of well-known heuristic functions

- ◇ Manhattan distance ( $L_1$ ): On a square grid that allows 4 directions of movement.
- ◇ Diagonal distance ( $L_\infty$ ): On a square grid that allows 8 directions of movement.
- ◇ Euclidean distance ( $L_2$ ): On a square grid that allows any direction of movement.


**Manhattan Distance**



**Diagonal Distance**



**Euclidean Distance**



# Admissible heuristics

E.g., for the 8-puzzle:

$h_1(n)$  = number of misplaced tiles

$h_2(n)$  = total Manhattan distance

(i.e., no. of squares from desired location of each tile)

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

$$\underline{\underline{h_1(S) = ??}}$$

$$\underline{\underline{h_2(S) = ??}}$$

# Admissible heuristics

E.g., for the 8-puzzle:

$h_1(n)$  = number of misplaced tiles

$h_2(n)$  = total Manhattan distance

(i.e., no. of squares from desired location of each tile)

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

$$h_1(S) = ?? \quad 7$$

$$h_2(S) = ?? \quad 2+3+3+2+4+2+0+2 = 18$$

## Dominance

If  $h_2(n) \geq h_1(n)$  for all  $n$  (both admissible)  
then  $h_2$  *dominates*  $h_1$  and is better for search

Typical search costs:

$d = 14$  IDS = 3,473,941 nodes

$A^*(h_1) = 539$  nodes

$A^*(h_2) = 113$  nodes

$d = 24$  IDS = too many nodes

$A^*(h_1) = 39,135$  nodes

$A^*(h_2) = 1,641$  nodes

## Relaxed problems

Admissible heuristics can be derived from the *exact* solution cost of a *relaxed* version of the problem

If the rules of the 8-puzzle are relaxed so that a tile can move *anywhere*, then  $h_1(n)$  gives the shortest solution

If the rules are relaxed so that a tile can move to *any adjacent square*, then  $h_2(n)$  gives the shortest solution

# Iterative improvement algorithms

In many optimization problems, *path* is irrelevant;  
the goal state itself is the solution


Then state space = set of “complete” configurations;  
find *optimal* configuration, e.g., Travelling Salesperson Problem  
or, find configuration satisfying constraints, e.g., n-queens

In such cases, can use *iterative improvement* algorithms;  
keep a single “current” state, try to improve it

Constant space, suitable for online as well as offline search

## Example: Travelling Salesperson Problem


Find the shortest tour that visits each city exactly once



Relaxed problem: let path be *any* structure that connects all cities  
 $\implies$  use minimum spanning tree as heuristic for the TSP

## Example: $n$ -queens

Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal





# Hill-climbing (or gradient ascent/descent)


“Like climbing Everest in thick fog with amnesia”

```
function HILL-CLIMBING(problem) returns a solution state
  inputs: problem, a problem
  local variables: current, a node
                     next, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    next ← a highest-valued successor of current
    if VALUE[next] < VALUE[current] then return current
    current ← next
  end
```

## Hill-climbing contd.

Problem: depending on initial state, can get stuck on local maxima



## Summary

Heuristics help reduce search cost,  
however, finding an optimal solution is still difficult.

Greedy best-first search is not optimal, but can be efficient.

A\* search is complete and optimal, but is prohibitive in memory.

Hill-climbing methods operate on complete-state formulations,  
require less memory, but are not optimal.

Examples of skills expected:

- ◇ Demonstrate operation of search algorithms
- ◇ Discuss and evaluate the properties of search algorithms
- ◇ Derive and compare heuristics for a problem