

# COMP30027 Machine Learning

## Instance-based Learning

Semester 1, 2019

Jeremy Nicholson & Tim Baldwin & Karin Verspoor



THE UNIVERSITY OF  
MELBOURNE

© 2019 The University of Melbourne

# Lecture Outline

## ① Instance-based Learning

## ② Comparing things

Sets of descriptors

Similarity metrics

## ③ Nearest Neighbour classification

## Reminder: Instances

- The input to a machine learning system consists of:
  - **Instances:** the individual, independent examples of a concept

*also known as **exemplars***

- Each instance is described by  $n$  attribute-value pairs.
- Each instance also has a class label.

# ML Example: the *Cool/Cute* Classifier

- According to my Tim's 2 y.o. son:

Entity	Class	Entity	Class
self	cute	sports car	cool
self as baby	???	tiger	cool
big brother (4 y.o.)	cool	Hello Kitty	cute
big sister (6 y.o.)	cute	spoon	???
Mummy	cute	water	???

- What would we predict the class for the following to be:  
*train, koala, book on ML*

## Instance-based learning (IBL)?

- IBL algorithms are supervised learning algorithms; they learn from labelled examples.
- Requires labelled examples.
- Directly “learn-by-example”.
  - Input: instances.
  - Model: Some kind of function that maps instances to categories.

# Lecture Outline

## ① Instance-based Learning

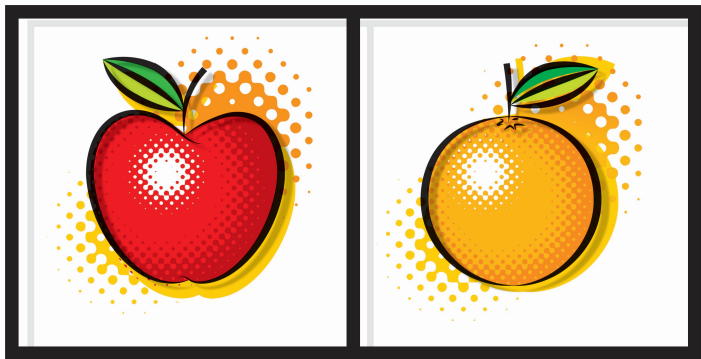
## ② Comparing things

- Sets of descriptors

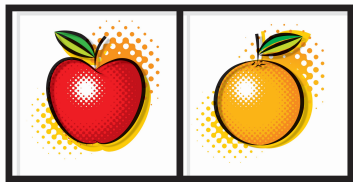
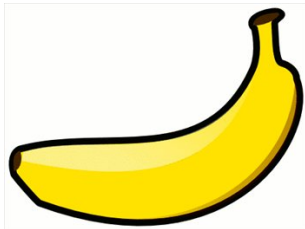
- Similarity metrics

## ③ Nearest Neighbour classification

# Compare and Contrast

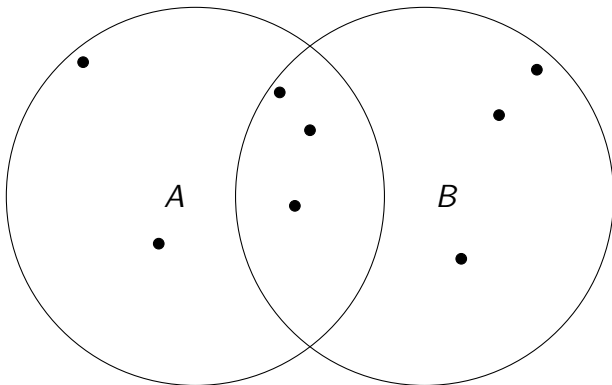


# Compare and Contrast





# Venn Diagram



## Similarity as Set intersection

Many similarity assessments can be framed as set intersection.

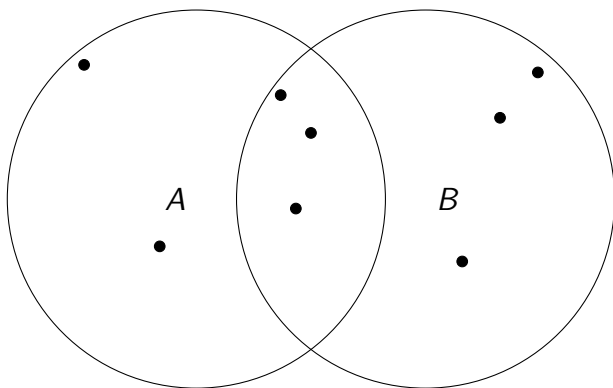
- Amazon: Book purchases
- Netflix: Movies that you have watched

### Refinements

- Rating sets (stars)
  - thresholding using ratings
  - different subsets for different ratings
- Categories of items
  - generalisation
  - book or movie genres

# Jaccard Similarity

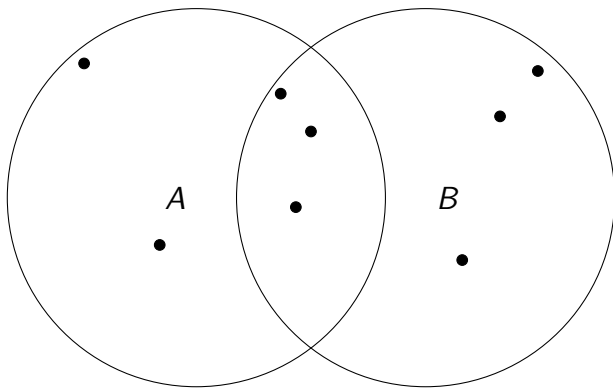
$$\frac{|A \cap B|}{|A \cup B|}$$



$$\text{sim}(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{3}{8}$$

# Dice Coefficient

$$\frac{2|A \cap B|}{|A| + |B|}$$



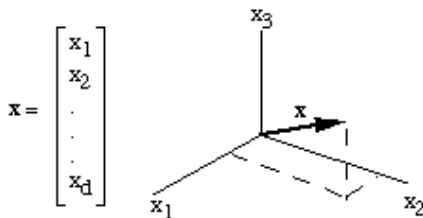
$$\text{sim}(A, B) = \frac{2|A \cap B|}{|A| + |B|} = \frac{2 * 3}{5 + 6} = \frac{6}{11}$$

## Feature vectors

- A *feature vector* is an n-dimensional vector of *features* that represent some object.
- A *feature* or *attribute* is any distinct aspect, quality, or characteristic of that object.
- Features may be nominal/symbolic/categorical/discrete (e.g. colour, gender)
- Features may be ordinal (e.g. cool < mild < hot [temperature])
- Features may be numeric/continuous (e.g., height, age)

## Feature vectors and vector space

A vector locates an instance (object, document, person, ...) as a point in an (orthogonal)  $n$ -space. The angle of the vector in that space is determined by the relative weight of each term.



- Similarity
  - Numerical measure of how alike two data objects are.
  - Is higher when objects are more alike.
  - Often falls in the range  $[0,1]$
- Dissimilarity
  - Numerical measure of how different are two data objects
  - Lower when objects are more alike
  - Minimum dissimilarity is often 0
  - Upper limit varies

# Similarity vs Distance

What is the relationship between similarity and distance?



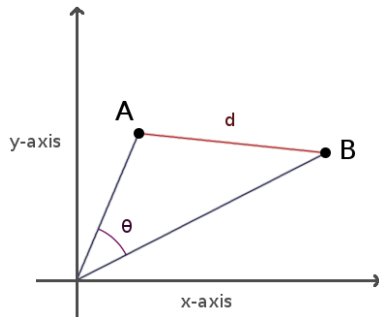
## Distance measures

A distance measure on a space is a function that takes two points in a space as arguments.

- No negative distances.  $d(x, y) \geq 0$
- Distances are positive, except for the distance from a point to itself.  $d(x, y) = 0$  if and only if  $x = y$
- Distance is symmetric.  $d(x, y) = d(y, x)$
- The *triangle inequality* typically holds.  
(Measures the length of *shortest path* between two points.)  
$$d(x, y) \leq d(x, z) + d(z, y)$$

# Euclidean Distance

Given two items  $A$  and  $B$ , and their feature vectors  $\mathbf{a}$  and  $\mathbf{b}$ , respectively, we can calculate their distance  $d$  in euclidean space:



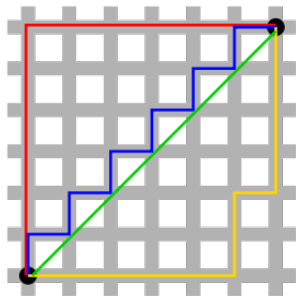
In n-dimensional space:

$$d(A, B) = \sqrt{\sum_{i=1}^n (a_i - b_i)^2}$$

# Manhattan Distance

[“City block” distance or “Taxicab geometry” or “ $L_1$  distance”]

Given two items  $A$  and  $B$ , and their corresponding feature vectors  $\mathbf{a}$  and  $\mathbf{b}$ , respectively, we can calculate their similarity via their distance  $d$  based on the absolute differences of their cartesian coordinates:

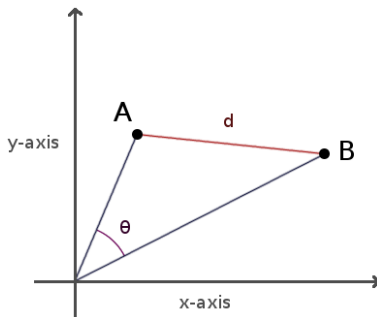


In  $n$ -dimensional space:

$$d(A, B) = \sum_{i=1}^n |a_i - b_i|$$

## Cosine Similarity

Given two items  $P$  and  $Q$ , and their feature vectors  $\mathbf{p}$  and  $\mathbf{q}$ , respectively, we can calculate their similarity via their **vector cosine** (the cosine of the angle  $\theta$  between the two vectors):



$$\cos(P, Q) = \frac{\mathbf{p} \cdot \mathbf{q}}{|\mathbf{p}| |\mathbf{q}|} = \frac{\sum_i p_i q_i}{\sqrt{\sum_i p_i^2} \sqrt{\sum_i q_i^2}}$$

# Lecture Outline

## ① Instance-based Learning

## ② Comparing things

Sets of descriptors

Similarity metrics

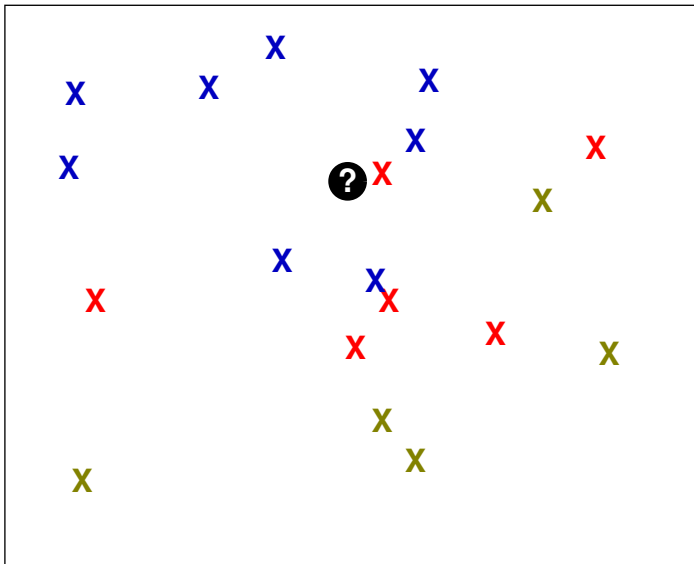
## ③ Nearest Neighbour classification

## What is a Nearest Neighbour?

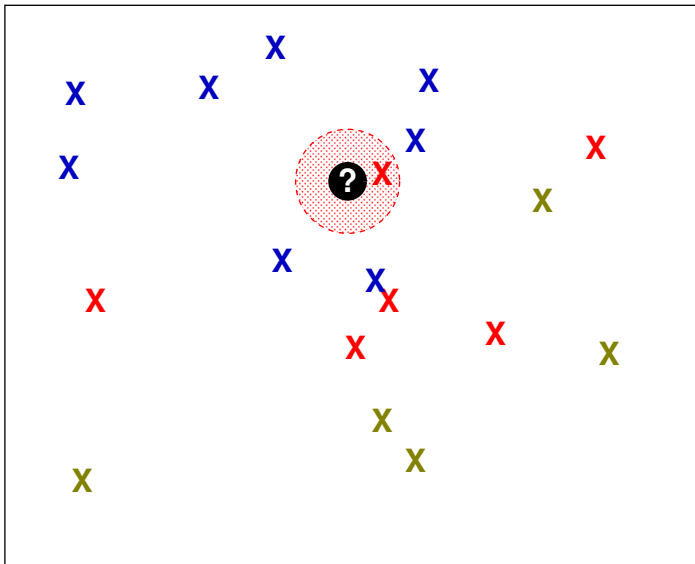
The closest point: maximum similarity or minimum distance.

$$d(x, y) = \min(d(x, z) | z \in Y)$$

# K neighbours

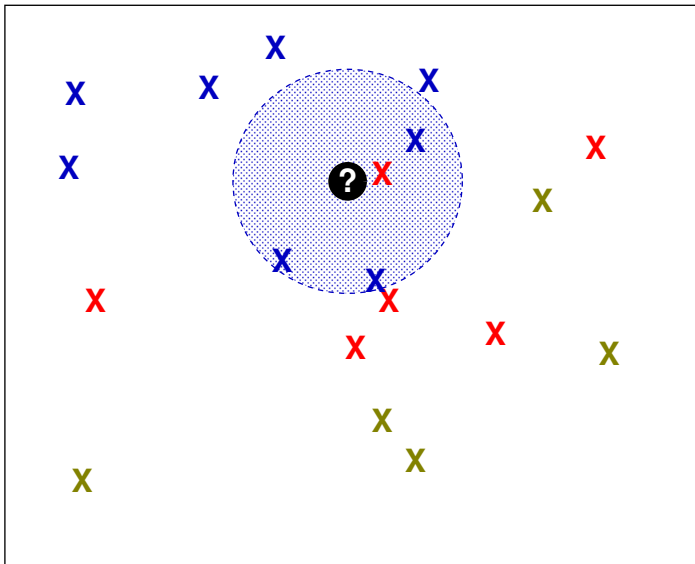


## K neighbours

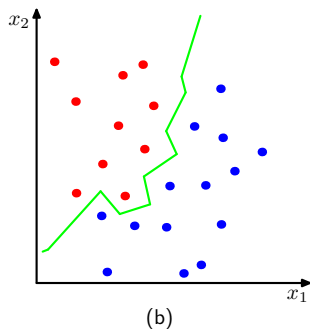
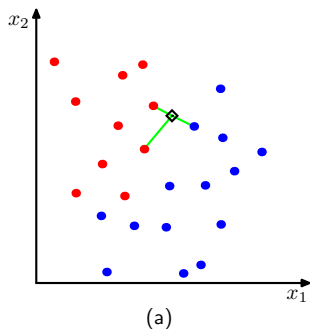




## K neighbours



# Nearest Neighbour methods in Classification



Given class assignments of existing data points, classify a new point (black).

- Consider the class membership of the  $K$  closest data points.
- For  $K = 1$ , the induced decision boundary. (b)

## Nearest Neighbour variants

**[1-NN]:** Classify the test input according to the class of the closest training instance.

**[ $K$ -NN]:** Classify the test input according to the majority class of the  $K$  nearest training instances.

**[weighted  $K$ -NN]:** Classify the test input according to the weighted accumulative class of the  $K$  nearest training instances, where weights are based on similarity of the input to each of the  $K$  neighbours.

**[offset-weighted  $K$ -NN]:** Classify the test input according to the weighted accumulative class of the  $K$  nearest training instances, where weights are based on similarity of the input to each of the  $K$  neighbours, factoring in an offset for the prior expectation of a test input being a member of that class.

## Weighting Strategies

- There are a number of strategies for weighting:
  - give each neighbour equal weight  
(= classify according to the **majority class** of set of neighbours)
  - weight the vote of each instance by its **inverse linear distance** from the test instance:

$$w_j = \begin{cases} \frac{d_k - d_j}{d_k - d_1} & \text{if } d_j \neq d_1 \\ 1 & \text{if } d_j = d_1 \end{cases}$$

where  $d_1$  is the nearest neighbour, and  $d_k$  is the furthest neighbour

- weight the vote of each instance by its **inverse distance** from the test instance:

$$w_j = \frac{1}{d_j + \epsilon}$$

# Voting Strategies in Action ( $k = 4$ )

- majority class voting:

$$\underline{yes} = 3 \text{ vs. } no = 1$$

Instance	Class	Distance
$d_1$	no	0
$d_2$	yes	1
$d_3$	yes	1.5
$d_4$	yes	2

- ILD-based voting:

$$yes = (\frac{1}{2} + \frac{1}{4} + 0)$$
$$\text{vs. } \underline{no} = 1$$

- ID-based voting ( $\epsilon = 0.5$ ):

$$yes = (\frac{1}{1.5} + \frac{1}{2} + \frac{1}{2.5})$$
$$\text{vs. } \underline{no} = \frac{1}{0.5}$$

# Breaking Ties

- In the case that we have an equal number of votes for a given class, we need some tie breaking mechanism:
  - random tie breaking
  - take class with highest prior probability
  - see if the addition of the  $k + 1$ th instance(s) breaks the tie

## Choosing the Value of $k$

- Smaller values of  $k$  tend to lead to lower classifier performance due to noise (overfitting)
- Larger values of  $k$  tend to drive the classifier performance toward Zero-R performance
- Generally trial and error over the training data is the only way of getting  $k$  just right

**Note:**  $k$  is generally set to an odd value ... why?

## Nearest Neighbour classification implementation I

A typical implementation involves the brute-force computation of distances between a test instance and every training instance.

- For  $N$  training instances in  $D$  dimensions, this approach scales as  $O(DN)$ .
- Efficient brute-force searches can be very competitive for small data samples.
- However, as the number of samples  $N$  grows, the brute-force approach quickly becomes infeasible.



# Nearest Neighbour classification implementation II

Why is  $k$ -Nearest Neighbour so slow?

- The **model** built by Naive Bayes/Decision Trees is generally much smaller than the dataset:
  - Predicting the class of a test instance requires approximately  $\mathcal{O}(CD)$  calculations for Naive Bayes, and  $\mathcal{O}(D)$  node traversals for a Decision Tree, given  $C$  classes and  $D$  attributes
- The **model** built by  $k$ -NN is the dataset itself:
  - $k$ -NN is *lazy*
  - The time we save in training is lost if we have to make many predictions

# Strengths and Weaknesses of NN methods

## Strengths

- Simple
- Can handle arbitrarily many classes
- Incremental (can add extra data to the classifier on the fly)

## Weaknesses

- We need a useful distance function.
- We need an averaging function for combining the labels of multiple training examples.
- Expensive (in terms of index accesses)
- Everything is done at run time (lazy learner)
- Prone to bias
- Arbitrary  $K$  value

# Summary

- Representing instances as vectors
- Measuring similarity
- What is  $k$ -Nearest Neighbour, and why do we call it an instance-based learning method?
- What parameters do we have to choose for  $k$ -NN?

## Readings:

- **Similarity:** Tan et al (2006), Section 2.4
- **NN classifier:** Tan et al (2006), Chapter 5, Section 5.2